

Assignment 1 Guide

BIOL4062/5062: Analysis of Biological Data

Reid Steele (Based on work by Ana Eguiguren)

2024-11-04

Table of contents

Introduction	4
Getting Help	4
Assignment Guidelines	6
General Advice	6
Submission Formatting	7
 I Assignment 1	 9
1 Assignment 1a: Principal Components Analysis	10
1.1 Looking at the Data	10
1.2 Transformations	15
1.3 Running PCA	22
1.3.1 Covariance Matrix	22
1.3.2 Correlation Matrix	29
1.3.3 Alternative Methods	32
1.4 Varimax Rotation (Optional)	38
1.5 Tips for your assignment:	40
 2 Assignment 1b: Linear Discriminant Analysis	 41
2.1 Looking at the data	41
2.2 MANOVA	44
2.3 Linear Discriminant Analysis	46
2.4 Model Selection	50
2.5 Plotting Probabilities	53
2.6 Tips for your assignment	56
 3 Assignment 1c: Cluster Analysis and Multidimensional Scaling	 57
3.1 Looking at the data	57
3.2 Calculating Dissimilarity	58
3.3 Hierarchical Cluster Analysis	59
3.3.1 Single linkage	59
3.3.2 Average Linkage	61
3.3.3 Complete Linkage	62
3.3.4 Ward Linkage	63

3.4	Multidimensional Scaling	64
3.4.1	Non-Metric MDS	64
3.4.2	Metric MDS	77
3.4.3	3D Plotting (Optional)	82
3.5	Mantel Test (Graduate Students Only)	83
3.6	Tips for your Assignment:	84
4	Assignment 1d: Multiple Linear Regression	86
4.1	Looking at the data	86
4.2	Considering Transformations	91
4.3	Simple Linear Regression	98
4.4	Multiple Linear Regression	107
4.5	Checking Assumptions	108
4.5.1	Independence	108
4.5.2	Linearity	110
4.5.3	Homoscedasticity	110
4.5.4	Normality	112
4.5.5	What if my assumptions aren't respected?	114
4.6	Model Selection	115
5.1	Tips for your Assignment:	124
6	Assignment 1e: Bayesian Data Analysis	125
6.1	Looking at the Data	125
6.2	Binomial GLM	126
6.3	Making it Bayesian	128
6.4	Adding Priors	136
6.5	Tips for your Assignment:	143

Introduction

Welcome to the assignment guide for BIOL4062/5062: Analysis of Biological Data.

This website is designed to walk you through the assignments for this class. It is a resource to help you figure out how to code for your assignments, and bring attention to key questions to ask yourself as you interpret your results, both statistically and biologically. Keep in mind that there are always different ways to get to the right answer with coding. The content here is not a monolith. You don't need to follow it if you don't want to (more in Assignment Guidelines), but make sure what you are doing is clear and sufficiently analogous to this guide, else you lose marks for being unclear, or running the wrong analyses.

Getting Help

Don't suffer in silence! If you need help on the assignments, there are multiple options available to you:

- Assignment Drop-In Sessions:
 - TA-run in-person help sessions the week before each assignment is due.
 - Optional, but recommended
 - * Even if you don't have questions, you may benefit from hearing the questions of others
- BrightSpace Discussion Board:
 - Feel free to ask questions on the BrightSpace discussion board, or peruse questions already asked
 - * You may find your answer without even asking!
 - Make sure to start your question with the assignment number
- Email
 - Feel free to ask questions, or set up an appointment

- Direct Assignment 1 questions to the TA, and Assignment 2 and/or class administration (e.g. extension requests) to the instructors

Without further ado, let's get into it!

Assignment Guidelines

This section will run through some general advice on how to put together your assignments, and how to make your TA's life easier marking them:

General Advice

1. Read the grading rubric!

- It is quite objective. There is little latitude for part marks if you are missing parts that you need.

2. You don't need to tell me the statistical theory or background (unless it's relevant to your answers)

- All you have to do is answer the questions in the assignment. Anything you write outside of that is just eating up your page limit.

3. Put your biological interpretations together at the end

- Your interpretations are more likely to make sense and easier to mark if you put them at the end, and include all of your results together in them rather than inserting them throughout piecemeal.
- **Also, make sure your biological interpretations are consistent with your data/results! They don't have to be correct, but they have to match your data.**

4. Make sure your figures are readable

- We can't tell if your interpretation of your figures is correct if I can't interpret your figures.
- All text on figures should be readable.
- If you use color, make sure that the colors you use are clearly distinguishable.

5. You're not alone!

- If you have questions, come to the drop-in sessions, read the discussion boards, or email the TA to ask questions or set up a meeting if those don't work for you.
- Do the first two even if you don't have questions: You may find the answer to questions you didn't know you had.
- You're also welcome to ask questions after an assignment if you want to know why you were graded the way you were, or if you have questions about the comments provided or what you may have done wrong.

6. Ask if you need an extension

- We're pretty reasonable.

Submission Formatting

Here are some guidelines on how to submit your assignment to make my life easier. **You won't lose marks for not following them, but I would greatly appreciate it if you did.**

1. Hand in your assignment in 3 parts:

a) Your assignment text

- All assignment 1s have a 2 page limit. Put all your text first, with figures and tables together separately at the end. It is easier for me to tell how many pages you wrote this way. You're not going to lose marks if you're slightly over (this is not a writing class), if you're going to lose marks for writing too many pages we're going to be able to tell anyways.
- Should be a doc or PDF file so it can be opened in BrightSpace. Doesn't matter if it's produced through word, markdown, etc, as long as it's in one of those two formats.

b) Your script, submitted as a .txt file

- Submitting as a .txt file allows us to open the script in BrightSpace rather than having to download it if it's a .R file. It is much more difficult to run your script should we need to if you paste it in your assignment document.

c) The data you read into your script

- This is just to make it easy to run your script if there is a mistake.

Note: If you do your assignments in markdown, you can combine a) and b)

2. Don't put your name on your assignments, in your scripts, or in any of your file names

- The BrightSpace system is anonymous so your assignments are marked blind. That doesn't work if you write your name.
- Delete your file paths in your script for submission if they have your name in them.
- **You do need to put your B0 number, data code number, and whether you're a graduate student or an undergraduate student.**

3. Follow the assignment guide

- You don't have to follow this guide to get full marks on the assignments (as always, there are many correct answers when it comes to coding). That said, it's easier to follow what you're doing if you're doing the same thing as everyone else.
- It's OK to do your own thing, but if you make a mistake, it's going to be much harder to help you out, and it's going to take significantly more effort to mark.

Part I

Assignment 1

1 Assignment 1a: Principal Components Analysis

Assignment 1a focuses on Principal Components Analysis (PCA). Think of PCA as a method of finding associations between data series.

For this tutorial, we're going to use the dataset in `fishcatch.csv`.

1.1 Looking at the Data

With any data analysis, step 1 is always to look at your data:

```
# Load in data
data = read.csv('fishcatch.csv')

# View data structure
head(data)
```


	Hauls	mackerel	bluefin	sardine	squid	limpet
1	1	1.851	55.60	0.058	6.00	0.0004
2	2	1.925	1.20	0.252	0.08	0.0027
3	3	2.506	1.56	0.133	0.06	0.0015
4	4	1.537	30.00	0.064	9.35	0.0013
5	5	1.795	0.04	0.086	4.70	0.0022
6	6	3.371	45.00	0.078	7.66	0.0006

```
dim(data)
```

```
[1] 25 6
```


Our data is a 25 row, 6 column data frame, describing catch of 5 different fisheries species (columns 2-6) caught across 25 hauls (column 1). We want to know if certain species are associated with each other. Lets look a little deeper at the data:

```
# Generate boxplots  
boxplot(data[, -1]) # Exclude haul
```




```
# look at data distribution  
# par(mfrow = c(3,2)) # 1 column 5 row grid plot  
hist(data$mackerel, breaks = 10)
```

Histogram of data\$mackerel




```
hist(data$bluefin, breaks = 10)
```

Histogram of data\$bluefin

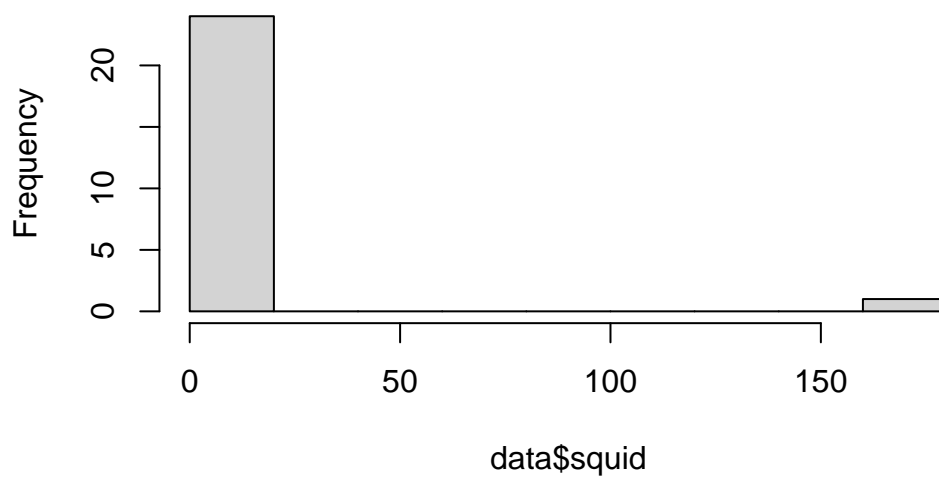


```
hist(data$sardine, breaks = 10)
```




```
hist(data$squid, breaks = 10)
```

Histogram of data\$squid



```
hist(data$limpet, breaks = 10)
```

Histogram of data\$limpet



A few things are immediately obvious from looking at our data:

1. There are some large outliers
2. The data scales vary greatly across species
3. The species all have relatively different distributions, none of which look normal.

Are these issues? How do we fix them?

1.2 Transformations

Look back at the PCA lecture. What are potential problems with PCA?

1. Covariance Matrix PCA requires data to be in the same units
2. Normality is desirable, but not essential
3. Precision is desirable, but not essential
4. Many zeroes in the data

We can fix issue 1 by logging our data:


```
# Create a new data object so we can log the data
data_log = data

# Log data
data_log[,-1] = log(data_log[,-1]) # Remember to exclude haul
```

Now that we've transformed the data, let's check for normality again:


```
# Generate histograms
# par(mfrow = c(3,2)) # 1 column 5 row grid plot
hist(data_log$mackerel, breaks = 10)
```

Histogram of data_log\$mackerel




```
hist(data_log$bluefin, breaks = 10)
```

Histogram of data_log\$bluefin

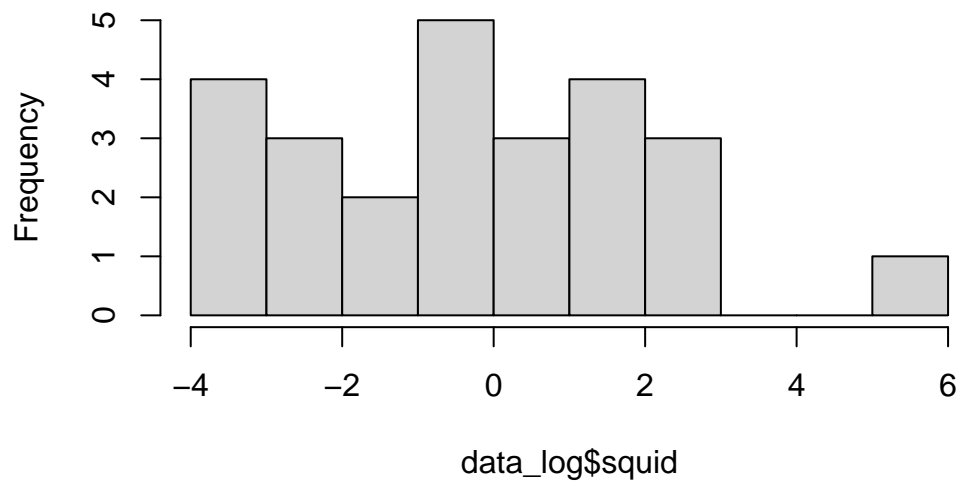



```
hist(data_log$sardine, breaks = 10)
```




```
hist(data_log$squid, breaks = 10)
```

Histogram of data_log\$squid



```
hist(data_log$limpet, breaks = 10)
```

Histogram of data_log\$limpet



These look much better. We can also confirm this statistically:

```
# Generate histograms  
shapiro.test(data_log$mackerel)
```

Shapiro-Wilk normality test

```
data: data_log$mackerel  
W = 0.9425, p-value = 0.1691
```

```
shapiro.test(data_log$bluefin)
```

Shapiro-Wilk normality test

```
data: data_log$bluefin  
W = 0.98186, p-value = 0.9193
```

```
shapiro.test(data_log$sardine)
```

Shapiro-Wilk normality test

```
data: data_log$sardine  
W = 0.94113, p-value = 0.1572
```

```
shapiro.test(data_log$squid)
```

Shapiro-Wilk normality test

```
data: data_log$squid  
W = 0.96226, p-value = 0.4613
```


```
shapiro.test(data_log$limpet)
```

Shapiro-Wilk normality test

```
data: data_log$limpet  
W = 0.96437, p-value = 0.5082
```

All 5 species fail to reject the null hypothesis that the data are normally distributed. Logging the data also helps deal with the outliers:

```
# Generate boxplots  
boxplot(data_log[, -1])
```



Note that we can only log the data if there are no zeroes:

```
# Generate test data  
data_test = data; data_test[1,6] = 0 # Change the first limpet value to 0  
  
# Try to log the data  
data_test[1,] # Print first row
```

```
Hauls mackerel bluefin sardine squid limpet  
1      1      1.851    55.6    0.058      6      0
```


```
log(data_test[,-1])[1,] # Print logs of the first row
```

```
mackerel bluefin sardine squid limpet
1 0.615726 4.018183 -2.847312 1.791759 -Inf
```

$\log(0)$ returns negative infinity. That's going to be a problem later in our analysis. We can fix that by adding a small increment before taking the log. Keep in mind though that each species has a different magnitude in this dataset, and adding an inappropriate increment could cause us trouble later:

```
# Test boxplots of different increments
boxplot(log(data_test$limpet), # Warning because of -Inf
        log(data_test$limpet + 1),
        log(data_test$limpet + 0.001),
        log(data_test$limpet + 0.000000001))
```

```
Warning in bplt(at[i], wid = width[i], stats = z$stats[, i], out =
z$out[z$group == : Outlier (-Inf) in boxplot 1 is not drawn
```



If the increment is too big, we eliminate the variance in our data. If the increment is too small, we create an outlier.

1.3 Running PCA

Now that we've checked and transformed our data, we're ready to run PCA. There are two kinds of PCA: We can run PCA on the Covariance Matrix, or the Correlation Matrix.

1.3.1 Covariance Matrix

We can run PCA on the covariance matrix as follows:


```
# Run PCA - Covariance
pca_1 = princomp(data_log[, -1]) # We don't want haul in our PCA!
summary(pca_1)
```

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
Standard deviation	2.8055354	1.3857803	1.3351790	0.55247629	0.182937780
Proportion of Variance	0.6607195	0.1612035	0.1496458	0.02562199	0.002809263
Cumulative Proportion	0.6607195	0.8219229	0.9715687	0.99719074	1.000000000

Running a summary on our PCA gives us the standard deviation of each principal component, the proportion of variance explained by each principal component, and the cumulative variance explained as we add each component. Here, we see the first principal component explains 66% of the variance. The second explains 16%, which adds up to 82% with the first component, and so on up to component 5. We can visualize the cumulative variance explained with a scree plot:

```
# Generate scree plot
plot(pca_1, type = 'l') # Scree is built into the plot for PCA
```



We see most of the variance is explained by component 1, then a similar lesser amount is explained by 2 and 3, followed by another drop to 4 and 5.

```
# Print loadings
print(loadings(pca_1),cutoff=0.00) #all loadings!
```

Loadings:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
mackerel	0.018	0.294	0.047	0.527	0.796
bluefin	-0.654	0.136	0.739	-0.089	-0.020
sardine	0.060	0.626	-0.015	0.520	-0.577
squid	-0.745	0.049	-0.664	0.029	0.018
limpet	0.116	0.707	-0.102	-0.665	0.183

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
SS loadings	1.0	1.0	1.0	1.0	1.0
Proportion Var	0.2	0.2	0.2	0.2	0.2
Cumulative Var	0.2	0.4	0.6	0.8	1.0

The PCA loadings are the correlations between the variables and each component. Here, we see bluefin and squid are strongly negatively correlated with component 1, while mackerel,

sardine, and limpet are weakly positively correlated with component 1. We can continue this type of interpretation through the other components as well.


Our PCA object also contains the PCA scores for each individual data point:

```
# Print PCA scores
head(pca_1$scores)
```

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
[1,]	-3.551007	-2.2507718	0.6725187	-0.1223707	-0.06754585
[2,]	2.484116	-0.6980669	0.4886052	-0.3932653	-0.53609582
[3,]	2.425206	-1.4149241	0.9556963	-0.2274184	-0.07560834
[4,]	-3.339469	-1.4723306	-0.2089590	-0.8854067	-0.03598265
[5,]	1.580988	-1.8002844	-4.6958830	-0.4313924	0.13590020
[6,]	-3.519487	-1.6187995	0.3362833	0.1040827	0.32134755

Scores are the value of each data point on each principal component. Lets try plotting them:


```
# Plot scores - components 1 and 2
plot(pca_1$scores[,1], # Scores on component 1
     pca_1$scores[,2], # Scores on component 3
     pch=16, # Point 16 (colored circle)
     xlab="1st principal component", ylab="2nd principal component", main="Scores plot") # A
```



This generates a scatterplot showing us the value of each data point in principal components 1 (x) and 2 (y). Now lets add on the loadings:

```
# Plot scores - components 1 and 2
plot(pca_1$scores[,1], # Scores on component 1
     pca_1$scores[,2], # Scores on component 3
     pch=16, # Point 16 (colored circle)
     xlab="1st principal component", ylab="2nd principal component", main="Scores plot") # A

# Add loadings to plot
arrows(0,0, # Draw arrows from zero
      pca_1$loadings[,1], # Draw to PC1 loading in X
      pca_1$loadings[,2], # Draw to PC2 loading in Y
      col="black", length = 0.1) # Arrow color and arrowhead length
text(pca_1$loadings[,1],pca_1$loadings[,2],names(data_log[,1]),cex=1.0 ,col="black") # Ad
```



The arrows are a little small, so let's add a scaling factor:


```
# Plot scores - components 1 and 2
plot(pca_1$scores[,1], # Scores on component 1
     pca_1$scores[,2], # Scores on component 3
```

```

    pch=16, # Point 16 (colored circle)
    xlab="1st principal component",ylab="2nd principal component",main="Scores plot") # A

# Add loadings to plot
sf = 3 # Scaling factor
sft = 3.2 # Scaling factor for text
arrows(0,0, # Draw arrows from zero
      pca_1$loadings[,1]*sf, # Draw to PC1 * scaling factor loading in X
      pca_1$loadings[,2]*sf, # Draw to PC2 * scaling factor loading in Y
      col="black", length = 0.1) # Arrow color and arrowhead length
text(pca_1$loadings[,1]*sft,pca_1$loadings[,2]*sft, names(data_log[,-1]), cex=1.0, col="bl

```



What about the haul number? Does that have an effect? Let's try adding that on as well:

```

# Create a color palette
colfunc = colorRampPalette(c('orangered1', 'turquoise2'))

# Plot scores - components 1 and 2
plot(pca_1$scores[,1], # Scores on component 1
     pca_1$scores[,2], # Scores on component 3
     pch=16, # Point 16 (colored circle)


```

```

col = colfunc(nrow(pca_1$scores)), # Color points by haul using our color palette
xlab="1st principal component",ylab="2nd principal component",main="Scores plot") # A

# Add loadings to plot
sf = 3 # Scaling factor
sft = 3.2 # Scaling factor for text
arrows(0,0, # Draw arrows from zero
      pca_1$loadings[,1]*sf, # Draw to PC1 * scaling factor loading in X
      pca_1$loadings[,2]*sf, # Draw to PC2 * scaling factor loading in Y
      col="black", length = 0.1) # Arrow color and arrowhead length
text(pca_1$loadings[,1]*sft,pca_1$loadings[,2]*sft, names(data_log[,-1]), cex=1.0, col="bl

```



Since we used color for haul, we need to add a legend:

```

# Set plot layout
layout(matrix(1:2,ncol=2), # 1 row, 2 columns
       width = c(2,1), # Width
       height = c(1,1)) # Height

# Create a color palette
colfunc = colorRampPalette(c('orangered1', 'turquoise2'))

```


```

# Plot scores - components 1 and 2
plot(pca_1$scores[,1], # Scores on component 1
     pca_1$scores[,2], # Scores on component 3
     pch=16, # Point 16 (colored circle)
     col = colfunc(nrow(pca_1$scores))), # Color points by haul using our color palette
     xlab="1st principal component", ylab="2nd principal component", main="Scores plot") # A

# Add loadings to plot
sf = 3 # Scaling factor
sft = 3.2 # Scaling factor for text
arrows(0,0, # Draw arrows from zero
       pca_1$loadings[,1]*sf, # Draw to PC1 * scaling factor loading in X
       pca_1$loadings[,2]*sf, # Draw to PC2 * scaling factor loading in Y
       col="black", length = 0.1) # Arrow color and arrowhead length
text(pca_1$loadings[,1]*sft, pca_1$loadings[,2]*sft, names(data_log[,-1]), cex=1.0, col="black")

# Generate legend
legend_image <- as.raster(matrix(colfunc(nrow(pca_1$scores))), ncol=1))
plot(c(0,2), c(0,1), type = 'n', axes = F, xlab = '', ylab = '', main = 'Haul')
text(x=1.5, y = seq(0,1,l=5), labels = seq(1,25,l=5))
rasterImage(legend_image, 0, 0, 1,1)

```



Now we have a completed scores plot with loadings arrows. How would you interpret this plot?

1.3.2 Correlation Matrix

Now let's try the correlation matrix. The correlation matrix performs the same analysis, but on standardized data. The `princomp()` function does this for us if we set `cor = T`:

```
# Run PCA - Correlation
pca_2 = princomp(data_log[, -1], cor = T)
summary(pca_2)
```

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
Standard deviation	1.595782	1.2503536	0.70708572	0.57220519	0.25041296
Proportion of Variance	0.509304	0.3126768	0.09999404	0.06548376	0.01254133
Cumulative Proportion	0.509304	0.8219809	0.92197491	0.98745867	1.00000000


```
# In case you don't believe me, heres the covariance matrix if we pre-standardize the data
pca_test = princomp(scale(data_log[, -1]))
summary(pca_test)
```

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
Standard deviation	1.563541	1.2250914	0.69279969	0.56064429	0.24535359
Proportion of Variance	0.509304	0.3126768	0.09999404	0.06548376	0.01254133
Cumulative Proportion	0.509304	0.8219809	0.92197491	0.98745867	1.00000000

Now we can go through the same pattern of analyses as we did for covariance:

```
# Generate scree plot
plot(pca_2, type = 'l') # Scree is built into the plot for PCA
```



```
# Print loadings
print(loadings(pca_2), cutoff=0.00) #all loadings!
```

Loadings:

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
mackerel	0.524	0.272	0.527	0.297	0.535
bluefin	-0.198	0.682	0.264	-0.651	-0.050
sardine	0.591	0.209	0.025	0.109	-0.771
squid	-0.233	0.645	-0.472	0.550	0.059
limpet	0.532	0.036	-0.655	-0.416	0.338

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5
SS loadings	1.0	1.0	1.0	1.0	1.0
Proportion Var	0.2	0.2	0.2	0.2	0.2
Cumulative Var	0.2	0.4	0.6	0.8	1.0

```
# Set plot layout
layout(matrix(1:2, ncol=2), # 1 row, 2 columns
       width = c(2,1), # Width
       height = c(1,1)) # Height
```

```


# Create a color palette
colfunc = colorRampPalette(c('orangered1', 'turquoise2'))

# Plot scores - components 1 and 2
plot(pca_2$scores[,1], # Scores on component 1
      pca_2$scores[,2], # Scores on component 3
      pch=16, # Point 16 (colored circle)
      col = colfunc(nrow(pca_2$scores)), # Color points by haul using our color palette
      xlab="1st principal component", ylab="2nd principal component", main="Scores plot") # A

# Add loadings to plot
sf = 3 # Scaling factor
sft = 3.2 # Scaling factor for text
arrows(0,0, # Draw arrows from zero
        pca_2$loadings[,1]*sf, # Draw to PC1 * scaling factor loading in X
        pca_2$loadings[,2]*sf, # Draw to PC2 * scaling factor loading in Y
        col="black", length = 0.1) # Arrow color and arrowhead length
text(pca_2$loadings[,1]*sft, pca_2$loadings[,2]*sft, names(data_log[, -1]), cex=1.0, col="black")

# Generate legend
legend_image <- as.raster(matrix(colfunc(nrow(pca_2$scores)), ncol=1))
plot(c(0,2), c(0,1), type = 'n', axes = F, xlab = '', ylab = '', main = 'Haul')
text(x=1.5, y = seq(0,1,l=5), labels = seq(1,25,l=5))
rasterImage(legend_image, 0, 0, 1,1)

```




How would you interpret this plot? Does it differ from the covariance plot?

1.3.3 Alternative Methods


There are a few other ways you can generate, and/or plot your PCAs if you prefer.

1.3.3.1 Biplot

```
# Exploring biplot
biplot(pca_1) # Covariance
```

```
biplot(pca_2) # Correlation
```



1.3.3.2 ggplot

```
library(ggplot2)

# ggplot version - Covariance

# turn PCA scores into data frame
pca_1_plot = data.frame(Haul = data_log$Haul, pca_1$scores)


# Turn PCA loadings into data frame (This gets a little complicated)
pca_1_loadings = as.data.frame(matrix(as.numeric(pca_1$loadings),
                                       dim(pca_1$loadings)[1], dim(pca_1$loadings)[2]))
colnames(pca_1_loadings) = colnames(pca_1_plot)[-1]

# Plot
ggplot(pca_1_plot, aes(x = Comp.1, y = Comp.2, color = Haul)) +

# Scores
geom_point() + scale_colour_distiller(palette = 15) +

# Loadings
geom_segment(data = pca_1_loadings, aes(x = 0, y = 0, xend = Comp.1, yend = Comp.2),
            arrow = arrow(length = unit(0.3, "cm"), type = "open", angle = 25),
            linewidth = 1, color = "darkblue") +

# Labels
geom_text(data = pca_1_loadings, color = 'darkblue', nudge_x = 0.2, nudge_y = 0.2, # Lab
          aes(x = Comp.1, y = Comp.2, label = colnames(data_log)[-1]))
```



```
# ggplot version - Correlation

# turn PCA scores into data frame
pca_2_plot = data.frame(Haul = data_log$Haul, pca_2$scores)

# Turn PCA loadings into data frame
pca_2_loadings = as.data.frame(matrix(as.numeric(pca_2$loadings),
                                         dim(pca_2$loadings)[1], dim(pca_2$loadings)[2]))
colnames(pca_2_loadings) = colnames(pca_2_plot)[-1]

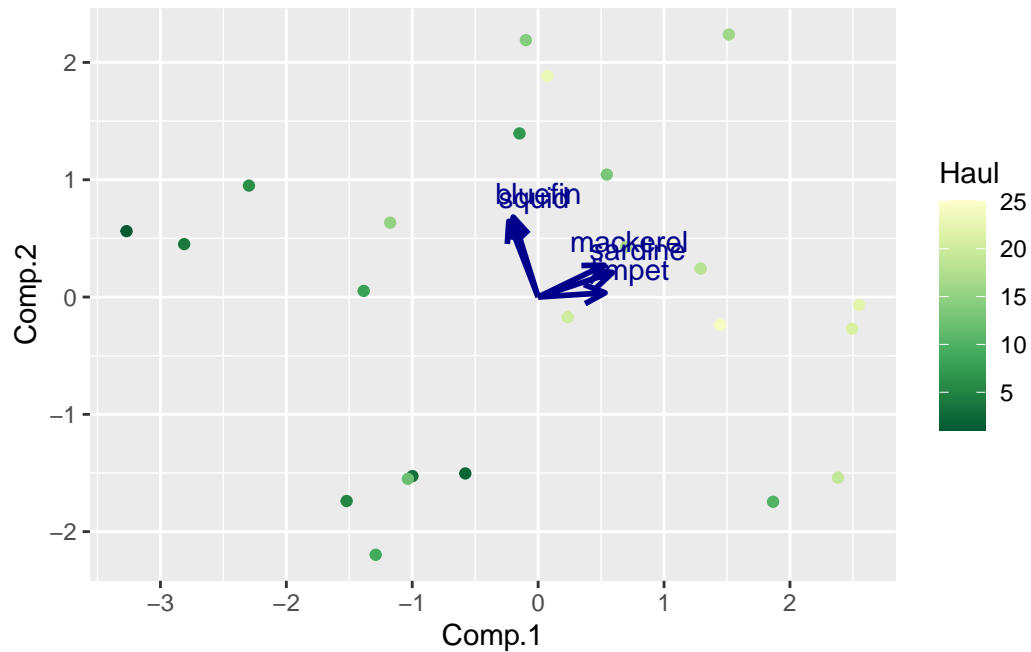
# Plot
ggplot(pca_2_plot, aes(x = Comp.1, y = Comp.2, color = Haul)) +

  # Scores
  geom_point() + scale_colour_distiller(palette = 15) +

  # Loadings
  geom_segment(data = pca_2_loadings, aes(x = 0, y = 0, xend = Comp.1, yend = Comp.2),
              arrow = arrow(length = unit(0.3, "cm"), type = "open", angle = 25),
              linewidth = 1, color = "darkblue") +

  # Labels
```

```
geom_text(data = pca_2_loadings, color = 'darkblue', nudge_x = 0.2, nudge_y = 0.2, # Lab
          aes(x = Comp.1, y = Comp.2, label = colnames(data_log)[-1]))
```




You can also run PCA using the `prcomp()` function instead of `princomp()`, setting `scale = T` if you want the correlation matrix. You can then use `autoplot()` with the `ggfortify` package to plot the results.

```
# ggplot v2
library(ggfortify)


# Run PCA - Covariance
pca_1a = prcomp(data_log[, -1])

# Run autoplot
autoplot(pca_1a, data = data_log, color = 'Hauls', loadings = T, loadings.label = T)
```



```
# Run PCA - Correlation
pca_2a = prcomp(data_log[, -1], scale = T)

# Run autoplot
autoplot(pca_2a, data = data_log, color = 'Hauls', loadings = T, loadings.label = T)
```



1.4 Varimax Rotation (Optional)


Varimax rotation attempts to improve the interpretability of PCA results by lining up loadings with the axes. This can be useful, particularly with large numbers of variables.

```
# Scaling factors
sf = 2.5
sft = 2.8

# Varimax rotation - Covariance
v1 = varimax(pca_1$loadings[,1:2])
v1_scores = pca_1$scores[,1:2] %>% v1$rotmat

# Plot scores - components 1 and 2
plot(v1_scores[,1], v1_scores[,2], pch=15, col = colfunc(nrow(v1_scores)),
     xlab="1st varimax component", ylab="2nd varimax component", main="varimax scores plot")


# Add loadings
arrows(0,0,v1$loadings[,1]*sf,v1$loadings[,2]*sf,col="black")
text(v1$loadings[,1]*sft,v1$loadings[,2]*sft,names(data_log[,-1]),asp=1,cex=1.0,col="black")
```



```
# Varimax rotation - Correlation
v2 = varimax(pca_2$loadings[,1:2])
v2_scores = pca_2$scores[,1:2] %*% v2$rotmat

# Plot scores - components 1 and 2
plot(v2_scores[,1], v2_scores[,2], pch=15, col = colfunc(nrow(v2_scores)),
     xlab="1st varimax component", ylab="2nd varimax component", main="varimax scores plot")

# Add loadings
arrows(0,0,v2$loadings[,1]*sf,v2$loadings[,2]*sf,col="black")
text(v2$loadings[,1]*sft,v2$loadings[,2]*sft,names(data_log[,-1]),asp=1,cex=1.0,col="black")
```



Note that it's pretty hard to tell the hauls apart using this color scale. Make sure your plots are always clear and readable.

1.5 Tips for your assignment:

Some things you may want to think about for your assignment:

1. Do your covariance and correlation plots differ? Do you think one is better suited to answering your research question? Why? Is your answer conceptual, or does it have to do with the results? Both?
2. How would you quantitatively examine the effect of haul on the PCA scores above? Is it associated with any of the principal components?
3. How would you interpret your statistical results biologically? You don't have to be right, but don't be vague, and don't contradict your results.

2 Assignment 1b: Linear Discriminant Analysis

Assignment 1b focuses on Linear Discriminant Analysis (LDA), also known as Canonical Variate Analysis. LDA is used to disclose relationships between groups, create models to differentiate between groups based on data, and discern the contribution of different variables to a model's ability to discriminate between groups.

For this tutorial, we'll be using `snake.csv`.

2.1 Looking at the data

```
# Load in data
snake = read.csv('snake.csv')

# Look at data
head(snake)
```

	Species	M1	M2	M3	M4	M5	M6
1	A	41.6	6.7	8.2	12.2	24.7	27.0
2	A	40.2	8.5	9.2	15.5	27.1	30.3
3	A	40.4	12.6	14.2	19.6	46.9	26.8
4	A	26.4	9.0	8.6	14.0	37.6	32.2
5	A	34.4	7.0	12.1	11.1	31.0	35.8
6	A	38.8	8.2	10.2	12.4	42.2	33.6

```
dim(snake)
```

```
[1] 35  7
```

Our data is a 35 row, 7 column data frame. The first column identifies the species of snake (A or B). The other columns are morphological measurements of each individual snake. We want to know if we can use the morphological measurements of the snakes to determine their species. Let's keep examining the data:

```

# Make a boxplot
library(tidyverse)

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2     3.4.4      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.0
v purrr       1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
  explicit

# Convert the data to long format so we can use ggplot
snake_long = pivot_longer(snake, # Enter data
                           colnames(snake)[-1], # Pivot all columns except species
                           names_to = 'Measurement', values_to = 'Value') # Feed labels to

# Lets take a look at the new data frame
head(snake_long)

# A tibble: 6 x 3
  Species Measurement Value
  <chr>      <chr>      <dbl>
1 " A " M1          41.6
2 " A " M2           6.7
3 " A " M3           8.2
4 " A " M4          12.2
5 " A " M5          24.7
6 " A " M6           27

# We've converted from wide format to long format,
# now all the data values are contained in a single column
# which is described by a metadata column

# You can also do this with melt from reshape2
library(reshape2)

```

Attaching package: 'reshape2'

The following object is masked from 'package:tidyr':


smiths

```
head(melt(snake))
```


Using Species as id variables

	Species	variable	value
1	A	M1	41.6
2	A	M1	40.2
3	A	M1	40.4
4	A	M1	26.4
5	A	M1	34.4
6	A	M1	38.8

```
# Let's make a boxplot  
ggplot(snake_long, aes(x = Measurement, y = Value, fill = Species)) +  
  geom_boxplot() + theme_classic()
```



```
# We can do this in R base plot too
boxplot(Value ~ Species*Measurement, # Plot value by species and measurement
        data = snake_long, col = c('coral', 'turquoise2'), # Color by species
        xaxt = 'n', xlab = 'Measurement') # Remove and label x axis
legend('topleft', legend = c('Species A', 'Species B'), fill = c('coral', 'turquoise2')) #
axis(1, at = seq(1.5,11.5,2), labels = colnames(snake)[-1]) # Add x axis back in with appr
```



Some of our measurements are very similar across species, and others are quite different. Do they differ statistically as a whole?

2.2 MANOVA

The purpose of LDA is to try to discriminate our snakes into species based on their measurements. However, that only makes sense to do if our two species of snake actually differ across the measurements. Our first step then is to discern whether our snake species differ as a multivariate whole. We'll do this using a MANOVA.

```
# Run MANOVA
sm = manova(cbind(M1,M2,M3,M4,M5,M6) ~ Species, data = snake)
summary(sm, test = 'Hotelling')
```

```

          Df Hotelling-Lawley approx F num Df den Df    Pr(>F)
Species    1          1.2263    5.7229      6    28 0.000552 ***
Residuals 33
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
summary(sm, test = 'Wilks')
```

```

          Df   Wilks approx F num Df den Df    Pr(>F)
Species    1 0.44917    5.7229      6    28 0.000552 ***
Residuals 33
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

By both the Hotelling's and Wilks' tests, our MANOVA is significant, indicating the snake species vary as a multivariate whole.

What about our assumptions though? Our MANOVA assumptions are normality, linearity, and homogeneity of covariances. You've been told to assume the latter, so let's skip that one.

```

# Testing normality
library(mvnormtest)
mshapiro.test(t(sm$residuals))

```

Shapiro-Wilk normality test

```

data:  Z
W = 0.91571, p-value = 0.01075


```

Uh oh, the residuals are significantly non-normal. Let's take a look at them visually:

```

# Residual histogram
hist(t(sm$residuals), breaks = 20)

```



Visually, our residuals actually look quite close to normal. There may be some slight skew, or outliers that are forcing our residuals to statistical non-normality. We might be able to fix this by removing multivariate outliers, or by transforming some of our data (feel free to play around with these ideas!), but based on the shape of our residuals, it is unlikely that our model is fatally biased, and we may end up doing more harm than good. Based on this, we can conclude that our two species have significantly different morphometries given the measurements provided.

2.3 Linear Discriminant Analysis

Now that we've confirmed our species differ as a multivariate whole, we can try to use LDA to build a model to predict which species each snake belongs to based on its measurements.

```
# LDA
library(MASS)
```

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

```
select
```

```
ldaf1 <- lda(Species ~ M1+M2+M3+M4+M5+M6, snake)
ldaf1
```

Call:

```
lda(Species ~ M1 + M2 + M3 + M4 + M5 + M6, data = snake)
```

Prior probabilities of groups:

	A	B
	0.2857143	0.7142857

Group means:

	M1	M2	M3	M4	M5	M6
A	32.700	9.410	10.16	15.54	35.290	28.950
B	31.496	12.128	10.18	16.78	47.356	21.752

Coefficients of linear discriminants:


	LD1
M1	0.01428023
M2	0.29104494
M3	-0.07327616
M4	-0.05544769
M5	0.03629586
M6	-0.17208517

Running our LDA object tells us the prior probabilities used for each species (the proportion of each species in the data), the group means for each measure on each species, and the linear discriminant (LD1) for each measure. We can then plot the LD1 value for each individual:

```
# Plot discriminant function analysis

# Create a data frame to plot
ldaf_plot = cbind(snake, # Data
                  predict(ldaf1)$x, # LD1 value for each individual given its measurements
                  index = seq(1,nrow(snake), 1)) # Row/Individual number

# Plot
plot(LD1 ~ index, data = ldaf_plot, col = as.factor(snake$Species), pch = 16)
legend('topleft', legend = c('A', 'B'), col = c(1, 2), pch = 16) # Add legend
```



Here we can see higher LD1 values are associated with species B, while lower LD1 values are associated with species A. This is just based on model fit however; how do we know we aren't overfitting? One way to avoid overfitting is by jackknifing (AKA leave-one-out cross validation in this context). This method runs the model once without each point in the dataset, then calculates the posterior probability that the left out point belongs to each species. Let's try it out:

```
# LDA 2, CV = T
ldaf2 = lda(Species ~ M1+M2+M3+M4+M5+M6, snake, CV = T)

# Gather posteriors
as.data.frame(cbind(ldaf2$posterior, # Pull posteriors from ldaf2
                    ResultantSpp=as.character(ldaf2$class))) # Pull predicted species (i.e
```

	A	B	ResultantSpp
1	0.897801237948675	0.102198762051325	A
2	0.957033498274347	0.0429665017256531	A
3	0.00486396795570833	0.995136032044292	B
4	0.939579607872302	0.0604203921276979	A
5	0.999020574119129	0.000979425880871094	A
6	0.958283942083953	0.041716057916047	A
7	0.859914694048175	0.140085305951825	A

8	0.0790276689479006	0.9209723310521	B
9	0.250711809994116	0.749288190005884	B
10	0.277233534989757	0.722766465010243	B
11	0.0654339037846645	0.934566096215335	B
12	8.13045175684633e-05	0.999918695482432	B
13	0.00857331675606214	0.991426683243938	B
14	0.119793120831736	0.880206879168264	B
15	0.868897347918874	0.131102652081126	A
16	0.291404395123413	0.708595604876587	B
17	0.580893601645515	0.419106398354485	A
18	0.407526292222816	0.592473707777184	B
19	0.097139347240766	0.902860652759234	B
20	0.0629455676122022	0.937054432387798	B
21	0.0262176442553781	0.973782355744622	B
22	0.0110464412594654	0.988953558740534	B
23	0.37967670676917	0.62032329323083	B
24	0.00300786068222619	0.996992139317774	B
25	0.0331152011340242	0.966884798865976	B
26	0.00270158005931187	0.997298419940688	B
27	0.0161164609849135	0.983883539015087	B
28	0.00346528534198866	0.996534714658011	B
29	0.761253716426843	0.238746283573157	A
30	0.0597294571669348	0.940270542833065	B
31	0.00139900114299064	0.998600998857009	B
32	0.0146304515487079	0.985369548451292	B
33	0.0215114427320867	0.978488557267913	B
34	0.00359029891803412	0.996409701081966	B
35	8.92739861715428e-05	0.999910726013828	B

How does this differ from the predictions from our first model?

```
# Pull ldaf1 model predictions
ldaf_pred = predict(ldaf1)$class

# Gather Predictions
ldaf_diff = data.frame(ldaf1 = as.character(ldaf_pred), ldaf2 = as.character(ldaf2$class))

# Add match column
ldaf_diff$match = (ldaf_diff$ldaf1 == ldaf_diff$ldaf2)

# Which ones are different?
ldaf_diff[which(ldaf_diff$match == F),]
```

	ldaf1	ldaf2	match
17	B	A	FALSE
29	B	A	FALSE

Individuals 17 and 29 both differed in species prediction between the model fit and the jackknife posterior probability. Now let's check the accuracy of our model fit:

```
# Calculate error
ldaf_wrong = length(which(ldaf_pred != snake$Species)) # Number of incorrect predictions
ldaf_err = ldaf_wrong/nrow(snake) # Divide by number of individuals for error

# Print error
ldaf_wrong
```

```
[1] 5
```

```
ldaf_err
```

```
[1] 0.1428571
```

Our model classified 5 out of 35 (~14.3%) of the snakes as the incorrect species, meaning 30/35 were correct (~85.7%). Not bad, but can we do better?

2.4 Model Selection

Our previous model used all 6 measurements, but do we really need all of them, or are some of them unhelpful (or even detrimental)? To test this, we can run model selection using the `stepclass()` function:

```
# stepclass package
library(klaR)

# Model selection (forward)
ms_f = stepclass(Species ~ M1+M2+M3+M4+M5+M6, data=snake,
                 method="lda", fold=35, direction="forward")
```

```
`stepwise classification', using 35-fold cross-validated correctness rate of method lda'.
```

35 observations of 6 variables in 2 classes; direction: forward

stop criterion: improvement less than 5%.

correctness rate: 0.85714; in: "M6"; variables (1): M6

hr.elapsed	min.elapsed	sec.elapsed
0.00	0.00	0.65

```
# Print model selection result
ms_f
```

```
method      : lda
final model  : Species ~ M6
<environment: 0x00000251e7c58178>
```

correctness rate = 0.8571

After model selection, we end up with a model using only M6 to predict species, with a correctness rate of 85.7%. This model has the same correctness as the full model, using only one measurement. In other words, this model is more **efficient** - it gets to the same accuracy using less information.

This model was generated using forward model selection, meaning the selection process works exclusively by adding variables to the model. We can also do the opposite:

```
# stepclass package
library(klaR)

# Model selection (backward)
ms_b = stepclass(Species ~ M1+M2+M3+M4+M5+M6, data=snae,
                 method="lda", fold=35, direction="backward")
```

`stepwise classification', using 35-fold cross-validated correctness rate of method lda'.

35 observations of 6 variables in 2 classes; direction: backward

stop criterion: improvement less than 5%.

```
correctness rate: 0.8; starting variables (6): M1, M2, M3, M4, M5, M6
correctness rate: 0.85714; out: "M5"; variables (5): M1, M2, M3, M4, M6
```

```
hr.elapsed min.elapsed sec.elapsed
      0.00      0.00      0.79
```

```
# Print model selection result
ms_b
```

```
method      : lda
final model  : Species ~ M1 + M2 + M3 + M4 + M6
<environment: 0x00000251e81bff40>
```

```
correctness rate = 0.8571
```

Backwards model selection works by removing variables from the full model. This means backwards selection should always return a model with a equal or more variables than forwards selection.

Lastly, we can run both:

```
# stepclass package
library(klaR)

# Model selection (forward)
ms_d = stepclass(Species ~ M1+M2+M3+M4+M5+M6,data=snae,
                 method="lda", fold=35, direction="both")
```

```
`stepwise classification', using 35-fold cross-validated correctness rate of method lda'.
```

```
35 observations of 6 variables in 2 classes; direction: both
```

```
stop criterion: improvement less than 5%.
```

```
correctness rate: 0.85714; in: "M6"; variables (1): M6
```

```
hr.elapsed min.elapsed sec.elapsed
      0.00      0.00      0.64
```

```
# Print model selection result
ms_d
```

```
method      : lda
final model  : Species ~ M6
<environment: 0x00000251e8a01b28>
```


```
correctness rate = 0.8571
```

2.5 Plotting Probabilities

Lets finish off by making some plots to visualize our LDA model results.

```
# Pick a model to plot
ldaf3 = lda(Species ~ M6, data = snake)



# Plot density curve
plot(ldaf3, dimen = 1, type = 'dens')
```



This plots the posterior probabilities of an individual belonging to either species given its LD1 value. Remember from earlier that species A is associated with lower LD1 values.



We can also make this plot as a histogram:

```
# Plot density curve
par(mar = c (4,4,4,4))
plot(ldaf3, dimen = 1, type = 'hist')
```



Or combine both plots:

```
# Plot density curve
par(mar = c (4,4,4,4))
plot(ldaf3, dimen = 1, type = 'both')
```




As always, we can also do this with ggplot too:

```
# Predict species
ldaf3_pred = predict(ldaf3)

# Plot
pred_species = as.data.frame(ldaf3_pred$x) # Gather LD1 values
pred_species$Species = snake$Species # Gather true species from data

# Plot
ggplot(pred_species, aes(x = LD1, fill = Species))+
  geom_density(alpha = 0.4)# alpha tells you how transparent the plots will be
```



2.6 Tips for your assignment

Some things you may want to think about for your assignment:

1. How would you pick which model you think is best? What factors would you consider? Are there any factors you would consider other than those discussed in this tutorial?
2. LDA also assumes the data are independent. Do we know this assumption is respected? Why or why not? What would constitute it not being respected?
3. How would you interpret your statistical results biologically (can be in terms of the snakes, how you would study them, or both)? You don't have to be right, but don't be vague, and don't contradict your results.

3 Assignment 1c: Cluster Analysis and Multidimensional Scaling

This assignment is centered on cluster analysis and multidimensional scaling (MDS), which are both methods of measuring associations within a group (e.g. associations between individuals within a population).

For this tutorial, we'll be using `monkey.csv`.

3.1 Looking at the data

You know the drill by now:

```
# Load in data
data = read.csv('monkey.csv', row.names = 1) # First column is row names
data # Print data
```

	ind1	ind2	ind3	ind4	ind5	ind6	ind7	ind8	ind9	ind10	ind11	ind12	ind13
ind1	21	2	2	10	2	2	8	0	0	8	14	12	4
ind2	2	21	16	2	16	8	2	2	4	4	4	0	2
ind3	2	16	21	0	10	16	2	0	2	4	4	0	2
ind4	10	2	0	21	2	2	16	2	2	8	12	8	4
ind5	2	16	10	2	21	10	2	4	0	2	4	4	2
ind6	2	8	16	2	10	21	4	2	0	0	0	4	4
ind7	8	2	2	16	2	4	21	4	2	16	8	8	4
ind8	0	2	0	2	4	2	4	21	0	2	0	0	0
ind9	0	4	2	2	0	0	2	0	21	0	4	0	0
ind10	8	4	4	8	2	0	16	2	0	21	14	14	2
ind11	14	4	4	12	4	0	8	0	4	14	21	12	4
ind12	12	0	0	8	4	4	8	0	0	14	12	21	2
ind13	4	2	2	4	2	4	4	0	0	2	4	2	21

Our data is a matrix containing the number of social interactions observed between individuals in a group of monkeys at the zoo. The matrix is symmetrical - the top/right half is identical to the bottom/left half.

3.2 Calculating Dissimilarity

For this assignment we'll be using 3 R functions: `hclust`, `metaMDS` (from the `vegan` package), `isoMDS` (from the `MASS` package), and `cmdscale()`. Let's see what type of input data those functions need:

```
# Check help functions
library(vegan)
library(MASS)
?hclust()
?metaMDS()
?isoMDS()
?cmdscale()
```

You'll notice all of these functions require a **dissimilarity matrix** produced by `dist`. Let's start by running `dist()`.

```
# Convert data to a dist object
dist = as.dist(data)
dist # Print dist
```

```
      ind1 ind2 ind3 ind4 ind5 ind6 ind7 ind8 ind9 ind10 ind11 ind12
ind2      2
ind3      2    16
ind4     10     2     0
ind5      2    16    10     2
ind6      2     8    16     2    10
ind7      8     2     2    16     2     4
ind8      0     2     0     2     4     2     4
ind9      0     4     2     2     0     0     2     0
ind10     8     4     4     8     2     0    16     2     0
ind11    14     4     4    12     4     0     8     0     4    14
ind12    12     0     0     8     4     4     8     0     0    14    12
ind13     4     2     2     4     2     4     4     0     0     2     4     2
```

Now our data is in a `dist` object. All of the redundant entries in the data have been removed.

Right now, our data reflects **similarity** (i.e. **high numbers reflect greater association between individuals**). We need to convert it to **dissimilarity**. Dissimilarity is simply the opposite of similarity. We can convert similarity to dissimilarity by subtracting each data value from the maximum of the data.

```
# Convert to dissimilarity
dist = max(dist) - dist
dist # Print dist
```

	ind1	ind2	ind3	ind4	ind5	ind6	ind7	ind8	ind9	ind10	ind11	ind12
ind2	14											
ind3	14	0										
ind4	6	14	16									
ind5	14	0	6	14								
ind6	14	8	0	14	6							
ind7	8	14	14	0	14	12						
ind8	16	14	16	14	12	14	12					
ind9	16	12	14	14	16	16	14	16				
ind10	8	12	12	8	14	16	0	14	16			
ind11	2	12	12	4	12	16	8	16	12	2		
ind12	4	16	16	8	12	12	8	16	16	2	4	
ind13	12	14	14	12	14	12	12	16	16	14	12	14

Now we're ready to run our analyses!

3.3 Hierarchical Cluster Analysis

Remember from lecture there are 4 types of hierarchical cluster analysis:

1. Single linkage
2. Average linkage
3. Complete linkage
4. Ward linkage

Let's run through them one by one:

3.3.1 Single linkage

We can run all 4 types of cluster analysis using the `hclust()` R function:

```
# run single linkage cluster analysis
clust_1 = hclust(dist, method = 'single')
clust_1 # print object
```

Call:


```
hclust(d = dist, method = "single")
```

Cluster method : single

Number of objects: 13

Printing the hclust object doesn't really tell us much. For more detail, we're going to have to plot it:

```
# Plot single linkage tree
plot(clust_1, hang = -1, main = 'Single linkage',
     ylab = 'Dissimilarity', # Label y axis
     xlab = '', sub = '') # Remove x-axis label
```



This outputs a tree showing the associations between our individual monkeys. dissimilarity is on the y-axis. The greater the distance between individuals on the y-axis, the greater their dissimilarity. Our tree has grouped the monkeys according to how frequently they interact with each other. For example. individuals 2, 3, 5, and 6 interact often, as evidenced by their low dissimilarity.

But how well does this tree fit the data? To answer that question, we need to calculate the cophenetic correlation coefficient (CCC):

```
# Calculate CCC
coph_1 = cophenetic(clust_1) # Get cophenetic
ccc_1 = cor(coph_1, dist) # Calculate correlation of the cophenetic with the data
ccc_1 # Print CCC
```


```
[1] 0.9036043
```

That's a pretty high correlation coefficient, indicating our dendrogram represented the structure in the original data very well. Let's try some other methods:

3.3.2 Average Linkage

```
# run cluster analysis
clust_2 = hclust(dist, method = 'average')

# Plot
plot(clust_2, hang = -1, main = 'Average linkage', ylab = 'Dissimilarity', xlab = '', sub
```




```
# Calculate CCC
coph_2 = cophenetic(clust_2)
ccc_2 = cor(coph_2, dist)
ccc_2
```

```
[1] 0.9288949
```

3.3.3 Complete Linkage

```
# run cluster analysis
clust_3 = hclust(dist, method = 'complete')

# Plot
plot(clust_3, hang = -1, main = 'Complete linkage', ylab = 'Dissimilarity', xlab = '', sub
```




```
# Calculate CCC
coph_3 = cophenetic(clust_3)
ccc_3 = cor(coph_3, dist)
ccc_3
```

```
[1] 0.9141956
```

3.3.4 Ward Linkage

```
# run cluster analysis
clust_4 = hclust(dist, method = 'ward.D')

# Plot
plot(clust_4, hang = -1, main = 'Ward linkage', ylab = 'Dissimilarity', xlab = '', sub = '')
```



```
# Calculate CCC
coph_4 = cophenetic(clust_4)
ccc_4 = cor(coph_4, dist)
ccc_4
```

```
[1] 0.7633159
```

Each method gives a slightly different tree and CCC value. Where are they similar? Where do they differ? Which one(s) would you trust? Why?

3.4 Multidimensional Scaling

Another method we can use to test for associations between our monkeys is multidimensional scaling (MDS). There are two types of MDS: non-metric, and metric MDS. Let's start with non-metric MDS.

3.4.1 Non-Metric MDS

```
# Run non-metric MDS - metaMDS
mds1 = metaMDS(dist, wascores = F)
```

```
Run 0 stress 0.07592385
Run 1 stress 0.07351184
... New best solution
... Procrustes: rmse 0.2259616  max resid 0.6667882
Run 2 stress 0.072393
... New best solution
... Procrustes: rmse 0.01792754  max resid 0.04236392
Run 3 stress 0.07875788
Run 4 stress 0.07239304
... Procrustes: rmse 9.65039e-05  max resid 0.0002134526
... Similar to previous best
Run 5 stress 0.07358653
Run 6 stress 0.07239303
... Procrustes: rmse 7.384877e-05  max resid 0.0001623498
... Similar to previous best
Run 7 stress 0.07358653
Run 8 stress 0.1787816
Run 9 stress 0.07239302
... Procrustes: rmse 6.496204e-05  max resid 0.0001435265
... Similar to previous best
Run 10 stress 0.072393
... Procrustes: rmse 5.516079e-06  max resid 1.019145e-05
... Similar to previous best
Run 11 stress 0.07575137
Run 12 stress 0.07358653
Run 13 stress 0.08055013
Run 14 stress 0.07366297
Run 15 stress 0.08566789
Run 16 stress 0.07366297
Run 17 stress 0.1183846
```



```
Run 18 stress 0.07358653
Run 19 stress 0.072393
... Procrustes: rmse 6.245426e-05  max resid 0.0001325748
... Similar to previous best
Run 20 stress 0.07366297
*** Best solution repeated 5 times
```

```
# Print mds results
mds1
```

```
Call:
metaMDS(comm = dist, wascores = F)

global Multidimensional Scaling using monoMDS

Data:      dist
Distance: user supplied

Dimensions: 2
Stress:     0.072393
Stress type 1, weak ties
Best solution was repeated 5 times in 20 tries
The best solution was from try 2 (random start)
Scaling: centring, PC rotation
Species: scores missing
```

By default, metaMDS has two dimensions. This MDS has a stress value of 0.072. Remember from lecture that stress < 0.10 is a “good representation”, so this MDS result is pretty good. If we want, we can test different numbers of dimensions (k) and create a scree plot to find the best one:

```
# Create a container object
scree = data.frame(k = 1:5, stress = NA)

# Loop through k 1 to 5
for(k in 1:5){

  # Run MDS
  mds = metaMDS(dist, wascores = F, k = k) # Set k to our loop index
```

```

# Pull out stress
scree[k,'stress'] = mds$stress # Fill kth row of the column 'stress' in scree

} # End loop

```

```

# Print results
scree

```

```


k      stress
1 1 2.360764e-01
2 2 7.239299e-02
3 3 9.147337e-05
4 4 9.314521e-05
5 5 9.724664e-05

```

```

# Make scree plot
plot(stress ~ k, data = scree, # Plot stress against k
      type = 'b', # Lines and points
      pch = 16) # Point 16 (filled circle)
abline(h = 0.1, lty = 'dashed') # Plot a dashed line at 0.1

```




We have an elbow at $k=3$, but we also get warnings that our dataset may be too small using $k=3$. The stress at $k=2$ is low enough that we can stick to using that.

Let's plot our results:

```
# Plot result
plot(mds1, type = 't')
```

species scores not available



Here we've plotted the values of our two MDS dimensions against each other for each individual. Similar to the cluster analysis, we see certain individuals are grouped together. Is it the same groups of individuals? What does that tell you about your results?

Let's try a different non-metric MDS function:

```
# Run non-metric MDS - isoMDS
mds2 = isoMDS(dist)
```

Error in isoMDS(dist): zero or negative distance between objects 2 and 3

Uh oh. This function doesn't like zeroes in the data. Let's fix that by translating our data to proportions, and adding a small increment.

```
# Translate to proportions
dist2 = dist/max(dist)

# Add an increment
dist2 = dist2 + 0.0001

# Print new dist
dist2
```

```
      ind1  ind2  ind3  ind4  ind5  ind6  ind7  ind8  ind9  ind10
ind2 0.8751
ind3 0.8751 0.0001
ind4 0.3751 0.8751 1.0001
ind5 0.8751 0.0001 0.3751 0.8751
ind6 0.8751 0.5001 0.0001 0.8751 0.3751
ind7 0.5001 0.8751 0.8751 0.0001 0.8751 0.7501
ind8 1.0001 0.8751 1.0001 0.8751 0.7501 0.8751 0.7501
ind9 1.0001 0.7501 0.8751 0.8751 1.0001 1.0001 0.8751 1.0001
ind10 0.5001 0.7501 0.7501 0.5001 0.8751 1.0001 0.0001 0.8751 1.0001
ind11 0.1251 0.7501 0.7501 0.2501 0.7501 1.0001 0.5001 1.0001 0.7501 0.1251
ind12 0.2501 1.0001 1.0001 0.5001 0.7501 0.7501 0.5001 1.0001 1.0001 0.1251
ind13 0.7501 0.8751 0.8751 0.7501 0.8751 0.7501 0.7501 1.0001 1.0001 0.8751
      ind11  ind12
ind2
ind3
ind4
ind5
ind6
ind7
ind8
ind9
ind10
ind11
ind12 0.2501
ind13 0.7501 0.8751
```

Let's make sure this doesn't mess with our results:

```

# Run non-metric MDS - metaMDS
mds1 = metaMDS(dist2, wascores = F)

Run 0 stress 0.07575137
Run 1 stress 0.07875788
Run 2 stress 0.07366297
... New best solution
... Procrustes: rmse 0.1934871  max resid 0.5648735
Run 3 stress 0.07366297
... Procrustes: rmse 2.166431e-06  max resid 3.493871e-06
... Similar to previous best
Run 4 stress 0.120356
Run 5 stress 0.08055013
Run 6 stress 0.07366297
... New best solution
... Procrustes: rmse 3.187788e-06  max resid 6.11983e-06
... Similar to previous best
Run 7 stress 0.07239302
... New best solution
... Procrustes: rmse 0.1707844  max resid 0.5155215
Run 8 stress 0.08571329
Run 9 stress 0.07239301
... New best solution
... Procrustes: rmse 1.713652e-05  max resid 3.265904e-05
... Similar to previous best
Run 10 stress 0.07875788
Run 11 stress 0.07366297
Run 12 stress 0.07239303
... Procrustes: rmse 3.940795e-05  max resid 8.67209e-05
... Similar to previous best
Run 13 stress 0.07227846
... New best solution
... Procrustes: rmse 0.01245006  max resid 0.03074824
Run 14 stress 0.07588698
Run 15 stress 0.07366297
Run 16 stress 0.07239304
... Procrustes: rmse 0.01248184  max resid 0.03099514
Run 17 stress 0.072393
... Procrustes: rmse 0.01248108  max resid 0.03099349
Run 18 stress 0.1180402
Run 19 stress 0.1183846
Run 20 stress 0.07366297

```

```
*** Best solution was not repeated -- monoMDS stopping criteria:  
10: stress ratio > sratmax  
10: scale factor of the gradient < sfgrmin
```

```
# Print mds results  
mds1
```

```
Call:  
metaMDS(comm = dist2, wascores = F)
```


```
global Multidimensional Scaling using monoMDS
```

```
Data:      dist2  
Distance: user supplied
```

```
Dimensions: 2  
Stress:      0.07227846  
Stress type 1, weak ties  
Best solution was not repeated after 20 tries  
The best solution was from try 13 (random start)  
Scaling: centring, PC rotation  
Species: scores missing
```

```
# Plot result  
plot(mds1, type = 't')
```

```
species scores not available
```



The values have shifted around a bit but the structure and interpretation of the plot is the same. Let's continue on:

```
# Run non-metric MDS - isoMDS
mds2 = isoMDS(dist2)
```

```
initial value 24.760322
iter 5 value 14.153502
iter 10 value 12.254154
iter 15 value 11.639473
iter 20 value 11.360460
final value 11.341572
converged
```

```
# Print output
mds2
```

```
$points
      [,1]      [,2]
ind1 0.5060798 0.103253718
ind2 -0.6157097 -0.285522725
```

```

ind3  -0.6133549 -0.310223762
ind4   0.5008471  0.144443835
ind5  -0.6264450 -0.279477605
ind6  -0.6228134 -0.325577873
ind7   0.4940123  0.147103149
ind8  -0.6783876  0.680257989
ind9  -0.2575043  1.075686777
ind10  0.5482152  0.033286470
ind11  0.5478840  0.082682890
ind12  0.5895070  0.001005794
ind13  0.2276697 -1.066918657

```

```

$stress
[1] 11.34157

```

The modelling algorithms seems to be a little different, and we end up with a different stress result - in this case, one that is above the 10% threshold (note that stress is in % in this function, unlike metaMDS where it is in proportion). Let's try another scree plot:

```

# Create a container object
scree = data.frame(k = 1:5, stress = NA)

# Loop through k 1 to 5
for(k in 1:5){

  # Run MDS
  mds = isoMDS(dist2, k = k) # Set k to our loop index

  # Pull out stress
  scree[k,'stress'] = mds$stress # Fill kth row of the column 'stress' in scree

} # End loop

# Print results
scree

```


```

   k    stress
1 1 36.54857293
2 2 11.34157190
3 3  0.04614441
4 4  0.12630298
5 5  0.19439990

```



```
# Make scree plot
plot(stress ~ k, data = scree, # Plot stress against k
     type = 'b', # Lines and points
     pch = 16) # Point 16 (filled circle)
abline(h = 10, lty = 'dashed') # Plot a dashed line at 0.1
```



In this case, it seems we're better off using 3 dimensions:

```
# Run non-metric MDS - isoMDS
mds2 = isoMDS(dist2, k = 3)
```

```
initial value 18.960422
iter 5 value 11.725940
iter 10 value 6.417141
iter 15 value 4.149185
iter 20 value 1.466748
iter 25 value 0.764657
iter 30 value 0.449114
iter 35 value 0.302911
iter 40 value 0.156116
iter 45 value 0.087536
```

```
iter 50 value 0.046144
final value 0.046144
stopped after 50 iterations
```

```
# Print output
mds2
```


```
$points
      [,1]      [,2]      [,3]
ind1  2.0028765  0.4079841 -0.4202376
ind2 -2.0901193 -1.3657278  1.2720375
ind3 -2.0910832 -1.3666943  1.2729356
ind4  2.0066696  0.4042391 -0.4161698
ind5 -2.0896284 -1.3631915  1.2769256
ind6 -2.0921681 -1.3637290  1.2724192
ind7  2.0032691  0.4111298 -0.4185160
ind8 -2.1600174  2.2033842 -1.8938700
ind9 -0.6641520  2.9972846  2.5707229
ind10 2.0009095  0.4046552 -0.4185896
ind11 2.0041166  0.4044773 -0.4197885
ind12 2.0019583  0.4023552 -0.4223419
ind13 -0.8326313 -2.1761669 -3.2555274
```

```
$stress
[1] 0.04614441
```

Let's plot our results:

```
# Plot isoMDS
plot(mds2$points[,1], mds2$points[,2], # MDS dimension 1 and 2 values
     type = 'n', # Don't plot any points
     xlab = 'Dim 1', ylab = 'Dim 2', main = 'Metric MDS') # Labelling

# Plot individual names
text(mds2$points[,1], mds2$points[,2], rownames(data))
```




All of our grouped individuals are plotted on top of each other. Let's try adding some random jiggle so we can see them

```
# Plot isoMDS
plot(mds2$points[,1], mds2$points[,2], # MDS dimension 1 and 2 values
     type = 'n', # Don't plot any points
     xlab = 'Dim 1', ylab = 'Dim 2', main = 'Metric MDS', # Labelling
     xlim = c(-2.5, 2.5), ylim = c(3, -3)) # Set axis limits

# Set random seed for consistency
set.seed(1212)

# Plot individual names
text(mds2$points[,1] + rnorm(13, 0, 0.2), # Add random values pulled from a
     mds2$points[,2] + rnorm(13, 0, 0.2), # normal distribution with mean 0, sd 0.2
     rownames(data)) # Add names
```



That's a bit better. We can also add some color to this plot if we want - say, individuals 6 to 9 are juveniles:

```
# Plot isoMDS
plot(mds2$points[,1], mds2$points[,2], # MDS dimension 1 and 2 values
     type = 'n', # Don't plot any points
     xlab = 'Dim 1', ylab = 'Dim 2', main = 'Metric MDS', # Labelling
     xlim = c(-2.5, 2.5), ylim = c(3, -3)) # Set axis limits
```

```
# Set random seed for consistency
set.seed(1212)
```

```
# Juvenile identifier
ad = c(rep(1,5), rep(0,4), rep(1,4)) # ad is 1 for first 5 and last 4
ad
```


```
[1] 1 1 1 1 1 0 0 0 0 1 1 1 1
```

```
# Plot individual names
text(mds2$points[,1] + rnorm(13, 0, 0.2), # Add random values pulled from a
     mds2$points[,2] + rnorm(13, 0, 0.2), # normal distribution with mean 0, sd 0.2
```

```

rownames(data), # Add names
col = ifelse(ad == 0, 'purple', 'orange')) # color
# Add a legend
legend('topright', legend = c('Juvenile', 'Adult'), fill = c('purple', 'orange'))

```



Does this plot match the previous one, and/or the cluster analyses?

3.4.2 Metric MDS

We can run metric MDS using the `cmdscale()` function:

```

# run metric MDS
mds3 = cmdscale(dist, eig = T)
mds3

```

```

$points
      [,1]      [,2]
ind1  5.84977914  2.7981654
ind2 -7.46899061 -0.4452479
ind3 -7.87360160  2.4742095

```

```

ind4    6.04970828 -1.1964346
ind5   -6.77887791  2.8518073
ind6   -7.21161499  3.9150940
ind7    4.79289905 -0.9896933
ind8   -2.46317365 -5.3304368
ind9   -1.86216019 -9.7770302
ind10   5.45394235  1.1317599
ind11   5.27120921 -0.1097836
ind12   6.20052885  3.6063451
ind13   0.04035206  1.0712450

$eig
 [1]  4.150450e+02  1.794715e+02  1.684147e+02  1.309366e+02  6.931537e+01
 [6]  5.811388e+01  3.306204e+01  1.780496e+01 -4.263256e-14 -8.643935e+00
[11] -2.208342e+01 -4.468321e+01 -7.952271e+01

$x
NULL

$ac
[1] 0

$GOF
[1] 0.4844901 0.5545014

```

Metric MDS doesn't have stress. Instead, we have to look at goodness of fit (GOF) to assess how well the analysis worked. GOF is similar to an R^2 value, where numbers closer to 1 indicate a better fit (though be wary of overfitting!). There are two different GOF values for each metric MDS.

As with the other MDS functions, `k` defaults to 2. We can make another scree plot:

```

# Create a container object
scree = data.frame(k = 1:5, GOF1 = NA, GOF2 = NA)

# Loop through k 1 to 5
for(k in 1:5){

  # Run MDS
  mds = cmdscale(dist, eig = T, k = k) # Set k to our loop index

  # Pull out stress

```

```

    scree[k,c(2,3)] = mds$GOF # Fill kth row of the GOF columns in scree


} # End loop

# Print results
scree

k      GOF1      GOF2
1 1 0.3382331 0.3871096
2 2 0.4844901 0.5545014
3 3 0.6217365 0.7115806
4 4 0.7284408 0.8337043
5 5 0.7849281 0.8983543

# Make scree plot
plot(GOF2 ~ k, data = scree, # Plot stress against k
     type = 'b', # Lines and points
     pch = 16, # Point 16 (filled circle)
     ylab = 'Goodness of Fit', ylim = c(0.3, 1))
points(GOF1 ~ k, data = scree, type = 'b', pch = 16, col = 'red') # Add second GOF value
abline(h = 0.1, lty = 'dashed') # Plot a dashed line at 0.1
legend('topleft', pch = 16, legend = c('GOF1', 'GOF2'), col = c('red', 'black')) # Add leg

```



Goodness of fit scales linearly, so what k to use is more of a judgement call.

```
# run metric MDS
mds3 = cmdscale(dist, k=4, eig = T)
mds3
```

```
$points
      [,1]      [,2]      [,3]      [,4]
ind1  5.84977914  2.7981654  1.2685787 -0.6738375
ind2 -7.46899061 -0.4452479  2.6496404  2.1091477
ind3 -7.87360160  2.4742095  3.4326421  1.2339175
ind4  6.04970828 -1.1964346 -0.9306802 -1.2319775
ind5 -6.77887791  2.8518073 -1.2465315  2.3989342
ind6 -7.21161499  3.9150940 -1.9699687 -2.4369605
ind7  4.79289905 -0.9896933 -2.7067201 -0.2043847
ind8 -2.46317365 -5.3304368 -9.4034890  2.0126521
ind9 -1.86216019 -9.7770302  5.2905629 -1.1532342
ind10  5.45394235  1.1317599  0.3565324  4.1657272
ind11  5.27120921 -0.1097836  4.1678975  1.4238728
ind12  6.20052885  3.6063451 -0.2995254  1.5298300
ind13  0.04035206  1.0712450 -0.6089393 -9.1736869
```



```

$eig
[1] 4.150450e+02 1.794715e+02 1.684147e+02 1.309366e+02 6.931537e+01
[6] 5.811388e+01 3.306204e+01 1.780496e+01 -4.263256e-14 -8.643935e+00
[11] -2.208342e+01 -4.468321e+01 -7.952271e+01

$x
NULL

$ac
[1] 0

$GOF
[1] 0.7284408 0.8337043

```

Let's plot the first two dimensions:

```


# Plot metric MDS
plot(mds3$points[,1], mds3$points[,2], # MDS dimension 1 and 2 values
     type = 'n', # Don't plot any points
     xlab = 'Dim 1', ylab = 'Dim 2', main = 'Metric MDS') # Labelling

# Plot individual names
text(mds3$points[,1], # Add random values pulled from a
     mds3$points[,2], # normal distribution with mean 0, sd 0.2
     rownames(data), # Add names
     col = ifelse(ad == 0, 'purple', 'orange')) # color

# Add a legend
legend('bottomright', legend = c('Juvenile', 'Adult'), fill = c('purple', 'orange'))

```

Metric MDS



3.4.3 3D Plotting (Optional)

It may not be necessary, but if your MDS has more than 2 dimensions, you can try plotting it in three dimensions and see if it helps:


```
library(plot3D)

# Prepare data to plot
x = mds3$points[,1]
y = mds3$points[,2]
z = mds3$points[,3]

# Create 3D plot
scatter3D(x,y,z, colvar = NULL, col = 'blue',
          pch = 16, cex = 0.5, bty = 'g', theta = 5)

# Add text
text3D(x,
       # Add some jiggle to the labels
       y+rnorm(13, mean = 0, sd = 0.5), z + rnorm(13, mean = 0, sd = 0.5),
       labels = names(mds3$points[,1]), add = T, colkey = F,
```

```
cex = 0.5, adj = 1, d = 2)
```



3.5 Mantel Test (Graduate Students Only)

We can infer to some extent whether juveniles and adults preferentially associate with each other from our colored MDS plots, but we can also test it statistically using a Mantel test. To run the Mantel test, we need to convert our adult index into a `dist` object:

```
# Create dist matrix for adults
ad_dist = dist(ad)
ad_dist
```

```

  1 2 3 4 5 6 7 8 9 10 11 12
2  0
3  0 0
4  0 0 0
5  0 0 0 0
6  1 1 1 1 1
7  1 1 1 1 1 0
8  1 1 1 1 1 0 0
```

```

9  1 1 1 1 1 0 0 0
10 0 0 0 0 0 1 1 1 1
11 0 0 0 0 0 1 1 1 1 0
12 0 0 0 0 0 1 1 1 1 0 0
13 0 0 0 0 0 1 1 1 1 0 0 0

```

Note this is dissimilarity: adult-juvenile pairs are assigned 1, and same-class pairs are assigned 0.

The Mantel test looks for correlation between this matrix and our original dissimilarity matrix, and statistically tests if the associations are different from what we would expect due to chance.

```

# Run mantel test
library(ade4)
mantel.rtest(ad_dist, dist, nrepet = 999)

```

```
Warning in is.euclid(m1): Zero distance(s)
```

```
Warning in is.euclid(m2): Zero distance(s)
```

```
Monte-Carlo test
```

```
Call: mantelnoneuclid(m1 = m1, m2 = m2, nrepet = nrepet)
```

```
Observation: 0.1686576
```

```
Based on 999 replicates
```

```
Simulated p-value: 0.073
```

```
Alternative hypothesis: greater
```

Std.Obs	Expectation	Variance
1.369210491	-0.001062026	0.015364686

It's very close, but we don't have statistically significant evidence that juveniles and adults associate preferentially with each other in this case.

3.6 Tips for your Assignment:

Some things you may want to think about for your assignment:

1. How would you pick which cluster analyses and MDS analyses are best for your data? Are they conceptual, or do they have to do with the results? Do they agree?
2. How would you interpret your statistical results biologically? You don't have to be right, but don't be vague, and don't contradict your results.

4 Assignment 1d: Multiple Linear Regression

This assignment is all about linear regression. Linear regression is used to model linear relationships between a dependent (response) variable and one or more independent (predictor) variables. Multiple linear regression involves multiple independent variables.

For this tutorial we're going to use `Schoenemann.csv`.

4.1 Looking at the data

```
# Read in data
data = read.csv('Schoenemann.csv')

# View data structure
head(data)
```

	Order	Family	Genus	Species	Location	Mass	Fat	FFWT	CNS
1	Carnivora	Felidae	Felis	canadensis	Alaska	7688.0	1120.0	6568.0	105.09
2	Carnivora	Felidae	Felis	rufus	Virginia	6152.0	738.0	5414.0	81.75
3	Carnivora	Mustelidae	Gulo	luscus	Alaska	9362.0	562.0	8800.0	85.36
4	Carnivora	Mustelidae	Mustela	erminea	Alaska	183.3	3.1	180.2	6.69
5	Carnivora	Mustelidae	Mustela	vison	Virginia	1032.0	66.0	966.0	18.06
6	Carnivora	Procyonidae	Procyon	lotor	Virginia	6040.0	1013.0	5027.0	58.31
	HEART	MUSCLE	BONE						
1	27.59	4341.45	631.18						
2	25.45	3600.31	552.23						
3	80.96	5271.20	879.12						
4	1.87	104.70	21.98						
5	7.63	581.53	80.27						
6	36.19	2920.69	517.78						

```
dim(data)
```

```
[1] 39 12
```

The Schoenemann dataset contains 39 observations of 12 variables, describing to the morphology of different species of mammals, along with metadata describing them. Let's start by getting rid of the metadata. We won't need it for this assignment.


```
# Remove metadata
data = data[,which(colnames(data) == 'Mass'):ncol(data)] # I do it this way to avoid hard

# check if it worked
head(data)
```



	Mass	Fat	FFWT	CNS	HEART	MUSCLE	BONE
1	7688.0	1120.0	6568.0	105.09	27.59	4341.45	631.18
2	6152.0	738.0	5414.0	81.75	25.45	3600.31	552.23
3	9362.0	562.0	8800.0	85.36	80.96	5271.20	879.12
4	183.3	3.1	180.2	6.69	1.87	104.70	21.98
5	1032.0	66.0	966.0	18.06	7.63	581.53	80.27
6	6040.0	1013.0	5027.0	58.31	36.19	2920.69	517.78

Now we only have numeric data left. Let's take a look at the data graphically.


```
# Looking at the data
boxplot(data)
```




```
# Loop through columns to create histograms
for(i in 1:ncol(data)){hist(data[,i], main = colnames(data)[i])} # Name histogram according
```




FFWT




CNS




HEART



MUSCLE





4.2 Considering Transformations


We can see that all these data are exponentially distributed - there is much more data at small values, and the data at higher values is spread out. If we run our regressions on this data, our assumptions are going to be violated:

```
# Run a test model and check assumptions
test = lm(CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE, data = data)

# Check for normality as an example


# Residual histogram
hist(residuals(test), 20)
```

Histogram of residuals(test)



```
# QQplot  
qqnorm(residuals(test))
```

Normal Q-Q Plot



```
# Statistical test for normality
shapiro.test(residuals(test))
```

Shapiro-Wilk normality test

```
data: residuals(test)
W = 0.87473, p-value = 0.0004494
```

Those diagnostics look... less than ideal.


This is a textbook case of when to apply a log transformation - remember logging is the opposite of exponentiating:

```
# Apply log transformation
data_1 = log(data)


# Check out the new data
head(data_1)
```



	Mass	Fat	FFWT	CNS	HEART	MUSCLE	BONE
1	8.947416	7.021084	8.789965	4.654817	3.3174534	8.375964	6.447591
2	8.724533	6.603944	8.596743	4.403666	3.2367157	8.188775	6.313965
3	9.144414	6.331502	9.082507	4.446878	4.3939552	8.570013	6.778921
4	5.211124	1.131402	5.194067	1.900614	0.6259384	4.651099	3.090133
5	6.939254	4.189655	6.873164	2.893700	2.0320878	6.365663	4.385396
6	8.706159	6.920672	8.522579	4.065774	3.5887828	7.979575	6.249550

```
# Looking at the data
boxplot(data_1)
```




```
# Loop through columns to create histograms
for(i in 1:ncol(data_1)){hist(data_1[,i], main = colnames(data_1)[i])} # Name histogram according to column name
```







CNS




HEART



MUSCLE



BONE



Now our data looks much more uniform.

Always remember that transforming your data incorrectly or unnecessarily can

do more harm than good. How do you decide if it is helpful to transform your data? What is the purpose of transforming your data? Think carefully about these questions for your assignment when you're deciding whether to transform the data for your assignment.

4.3 Simple Linear Regression

Now that our data is good to go, we're going to run some simple linear regressions to predict central nervous system mass (CNS). Simple linear regressions only have one predictor variable. Linear regression is run using the `lm()` command:

```
# Run simple linear regressions - Mass
m1 = lm(CNS ~ Mass, data = data_1) # run model
summary(m1) # model summary
```

Call:

```
lm(formula = CNS ~ Mass, data = data_1)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.77785	-0.20227	-0.05439	0.19607	0.78453

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.79097	0.14844	-18.80	<2e-16 ***
Mass	0.77105	0.02556	30.16	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3657 on 37 degrees of freedom

Multiple R-squared: 0.9609, Adjusted R-squared: 0.9599

F-statistic: 909.7 on 1 and 37 DF, p-value: < 2.2e-16

```
# Run simple linear regressions - Fat
m2 = lm(CNS ~ Fat, data = data_1) # run model
summary(m2) # model summary
```

Call:

```
lm(formula = CNS ~ Fat, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.22918	-0.41510	0.01431	0.36008	1.38000

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.15712	0.13223	-1.188	0.242
Fat	0.59903	0.03518	17.028	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6223 on 37 degrees of freedom

Multiple R-squared: 0.8868, Adjusted R-squared: 0.8838

F-statistic: 289.9 on 1 and 37 DF, p-value: < 2.2e-16

```
# Run simple linear regressions - FFWT
m3 = lm(CNS ~ FFWT, data = data_1) # run model
summary(m3) # model summary
```

Call:

```
lm(formula = CNS ~ FFWT, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.8057	-0.2112	-0.0535	0.1907	0.7654

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.80193	0.14382	-19.48	<2e-16 ***
FFWT	0.78494	0.02515	31.20	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3539 on 37 degrees of freedom

Multiple R-squared: 0.9634, Adjusted R-squared: 0.9624

F-statistic: 973.7 on 1 and 37 DF, p-value: < 2.2e-16

```
# Run simple linear regressions - HEART
m4 = lm(CNS ~ HEART, data = data_1) # run model
summary(m4) # model summary
```

Call:

```
lm(formula = CNS ~ HEART, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.75646	-0.16000	-0.03248	0.15018	0.85234

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.99514	0.05853	17.00	<2e-16 ***
HEART	0.88201	0.02872	30.71	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3594 on 37 degrees of freedom

Multiple R-squared: 0.9622, Adjusted R-squared: 0.9612

F-statistic: 943 on 1 and 37 DF, p-value: < 2.2e-16

```
# Run simple linear regressions - MUSCLE
m5 = lm(CNS ~ MUSCLE, data = data_1) # run model
summary(m5) # model summary
```

Call:

```
lm(formula = CNS ~ MUSCLE, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.82059	-0.15588	-0.00489	0.17331	0.80475

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.20579	0.11856	-18.61	<2e-16 ***
MUSCLE	0.76488	0.02296	33.31	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3323 on 37 degrees of freedom

Multiple R-squared: 0.9677, Adjusted R-squared: 0.9669

F-statistic: 1109 on 1 and 37 DF, p-value: < 2.2e-16

```
# Run simple linear regressions - BONE
m6 = lm(CNS ~ BONE, data = data_1) # run model
summary(m6) # model summary
```

Call:

```
lm(formula = CNS ~ BONE, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.10309	-0.24611	0.01155	0.25195	0.63931

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.0497	0.1042	-10.07	3.75e-12 ***
BONE	0.7856	0.0277	28.36	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3879 on 37 degrees of freedom

Multiple R-squared: 0.956, Adjusted R-squared: 0.9548


F-statistic: 804.4 on 1 and 37 DF, p-value: < 2.2e-16

In this case, it looks like all of our variables are strong, significant predictors with high R^2 values.


Let's plot all of these regressions:

```
# Plot regressions


# Plot simple linear regressions - Mass
plot(CNS ~ Mass, data = data_1, pch = 16) # plot points
abline(m1, lwd = 2, col = 'red') # Plot model
```




```
# Plot simple linear regressions - Fat
plot(CNS ~ Fat, data = data_1, pch = 16) # plot points
abline(m2, lwd = 2, col = 'red') # Plot model
```




```
# Plot simple linear regressions - FFWT
plot(CNS ~ FFWT, data = data_1, pch = 16) # plot points
abline(m3, lwd = 2, col = 'red') # Plot model
```




```
# Plot simple linear regressions - HEART
plot(CNS ~ HEART, data = data_1, pch = 16) # plot points
abline(m4, lwd = 2, col = 'red') # Plot model
```

```
# Plot simple linear regressions - MUSCLE  
plot(CNS ~ MUSCLE, data = data_1, pch = 16) # plot points  
abline(m5, lwd = 2, col = 'red') # Plot model
```



```
# Plot simple linear regressions - BONE
plot(CNS ~ BONE, data = data_1, pch = 16) # plot points
abline(m6, lwd = 2, col = 'red') # Plot model
```



All of the regression slopes are positive. This makes sense - larger animals tend to have larger brains.

4.4 Multiple Linear Regression

We've made 6 models using 1 variable. Now, let's try making 1 model with 6 variables:

```
# Run full model
m7 = lm(CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE, data = data_1)
summary(m7)
```

Call:

```
lm(formula = CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE,
    data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.72690	-0.12073	0.00376	0.08672	0.85638

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.61138	1.18496	-0.516	0.6094
Mass	-0.46867	2.14250	-0.219	0.8282
Fat	-0.06818	0.15489	-0.440	0.6628
FFWT	-0.02606	2.30347	-0.011	0.9910
HEART	0.41894	0.21913	1.912	0.0649
MUSCLE	1.03524	0.61123	1.694	0.1000
BONE	-0.06339	0.32054	-0.198	0.8445

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3286 on 32 degrees of freedom

Multiple R-squared: 0.9727, Adjusted R-squared: 0.9676

F-statistic: 190.1 on 6 and 32 DF, p-value: < 2.2e-16

Our full model has a very(!) high R^2 value, and contrary to the simple linear regressions where every predictor was significant, none of our predictors are considered significant in the final model at $\alpha = 0.05$. Why do you think that is?

4.5 Checking Assumptions

Now that we've run our full model, it's time to check its assumptions. Those assumptions are **Independence, Linearity, Homoscedasticity, and Normality**. By now, you should be familiar with what these all mean, but let's run through them anyways:

4.5.1 Independence

The assumption of independence states that the value of each data point ('datum', if you will) is independent of all other data points. Some of the ways in which it could be violated may not be testable (e.g. if they have to do with how the data was collected), but what we *can* test for is **autocorrelation**. Autocorrelation translates to self correlation (auto = self). We can test for autocorrelation statistically using a Durbin-Watson test, and visually using an autocorrelation function on the residuals:

```
library(lmtest)
```

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

```
# Durbin-watson test  
dwtest(m7)
```


Durbin-Watson test

data: m7

DW = 2.4315, p-value = 0.8545

alternative hypothesis: true autocorrelation is greater than 0

```
# Autocorrelation function  
acf(residuals(m7))
```



The Durbin-Watson test returns an insignificant p-value, indicating no autocorrelation structure is present. The ACF plots the correlation coefficient of the data against itself using lags. Lag 0 correlates the data against itself, which is always 1. Lag 1 correlates each data point against the point after it, and so on. All of the correlation coefficients are between the blue lines, so again, we have no autocorrelation structure, and we can say independence is respected.

4.5.2 Linearity

The assumption of linearity states that the response variable consistently scales linearly with its predictors. We can test for linearity statistically using Ramsey's RESET test on our model:

```
# Run RESET test
resettest(m7)
```

RESET test

```
data:  m7
RESET = 0.12784, df1 = 2, df2 = 30, p-value = 0.8805
```

In this case, the p-value is not significant, meaning the assumption of linearity is respected.

4.5.3 Homoscedasticity


The assumption of homoscedasticity is that the variance in the data is independent of the value of the data - i.e. the variance in the data is consistent. We can test this statistically using the Breusch-Pagan test, and visually by plotting the model residuals against the fitted values.

```
# Run Breusch-Pagan test
bptest(m7)
```


studentized Breusch-Pagan test

```
data:  m7
BP = 15.974, df = 6, p-value = 0.01389
```

```
# Plot residuals vs fitted
plot(m7$residuals ~ m7$fitted.values, pch = 16); abline(h = 0)
```



```
# Can also be done using plot.lm, ?plot.lm for details
plot(m7, 1)
```




The Breusch-Pagan returns a significant p-value, indicating the assumption of homoscedasticity is violated. We can see in the residuals versus fitted plot that the variance in the data is smaller at low values than it is at higher values (the points on the left of the plot are clustered more closely than they are on the right). Let's come back to this later.

4.5.4 Normality

The assumption of normality states that the residuals of our model should be normally distributed. If they aren't, that would indicate that our model is biased towards overprediction or underprediction in some way. As we did earlier in the transformation section, we can check for normality visually by looking at histograms and QQ plots of our residuals, and statistically by running a Shapiro-Wilk test on the residuals.


```
# Residual histogram
hist(residuals(m7))
```


Histogram of residuals(m7)




```
# QQplot  
qqnorm(residuals(m7))
```

Normal Q-Q Plot



```
# Can also use plot.lm for qqplot
plot(m7, 2)
```



```
# Statistical test for normality
shapiro.test(residuals(m7))
```

Shapiro-Wilk normality test

```
data: residuals(m7)
W = 0.95391, p-value = 0.1113
```

The Shapiro-Wilk test p-value is not significant (though it treads close), meaning the assumption of normality is respected. The residual histogram largely looks normal, and the QQ plot tails start to pull off the line at high and low values, possibly indicating outliers are causing us some trouble, but not enough to violate the assumption.

4.5.5 What if my assumptions aren't respected?

The typical fixes for violated assumptions are data transformations, and the removal of outliers. In our case, we pass all assumptions except for homoscedasticity. We've already transformed

our data to meet the assumption of normality, so further transformation is likely off the table, though we could potentially try different transformations. We could also try outlier removal - our model diagnostics using `plot.lm()` identify three outliers - 19, 20, and 28. Feel free to play around with removing outliers if you want.

Keep in mind that data transformations and removing outliers both represent trade-offs. Removing outliers may help meet your model assumptions, but you may also be removing data that reflects reality from your model. In that case, is it really helping you to remove outliers? Similarly, transforming your data may help you meet your assumptions, but in a case like this, transforming our data further or in a different way could end up violating other assumptions. Sometimes the best way to deal with violated assumptions is simply to state that they are violated and think about what that means for the interpretation of your model. Play around with all these different ideas, and come up with what you think is best. At the end of the day, a lot of statistical choices are judgement calls, with no perfect right answer.

4.6 Model Selection

In assignment 1b, we created 1 model with 6 variables, then tested if we could get a similarly effective model using fewer variables - i.e. a more **efficient** model. Let's do the same thing here:

```
# Stepwise model selection - forward
m8 = step(m7, direction = 'forward')
```

```
Start:  AIC=-80.52
CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE
```

```
summary(m8)
```

Call:

```
lm(formula = CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE,
    data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.72690	-0.12073	0.00376	0.08672	0.85638

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
--	----------	------------	---------	----------

(Intercept)	-0.61138	1.18496	-0.516	0.6094
Mass	-0.46867	2.14250	-0.219	0.8282
Fat	-0.06818	0.15489	-0.440	0.6628
FFWT	-0.02606	2.30347	-0.011	0.9910
HEART	0.41894	0.21913	1.912	0.0649 .
MUSCLE	1.03524	0.61123	1.694	0.1000
BONE	-0.06339	0.32054	-0.198	0.8445

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3286 on 32 degrees of freedom

Multiple R-squared: 0.9727, Adjusted R-squared: 0.9676

F-statistic: 190.1 on 6 and 32 DF, p-value: < 2.2e-16

Running forward model selection cuts the model down to 3 variables. As with 1b, we can also do backward:

```
# Stepwise model selection - backwards
m9 = step(m7, direction = 'backward')
```

Start: AIC=-80.52

CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE

	Df	Sum of Sq	RSS	AIC
- FFWT	1	0.00001	3.4552	-82.523
- BONE	1	0.00422	3.4594	-82.476
- Mass	1	0.00517	3.4604	-82.465
- Fat	1	0.02092	3.4761	-82.288
<none>			3.4552	-80.523
- MUSCLE	1	0.30975	3.7650	-79.175
- HEART	1	0.39468	3.8499	-78.305

Step: AIC=-82.52

CNS ~ Mass + Fat + HEART + MUSCLE + BONE

	Df	Sum of Sq	RSS	AIC
- BONE	1	0.00487	3.4601	-84.468
- Fat	1	0.04499	3.5002	-84.019
- Mass	1	0.05110	3.5063	-83.951
<none>			3.4552	-82.523
- MUSCLE	1	0.38148	3.8367	-80.439
- HEART	1	0.39621	3.8514	-80.289

```
Step: AIC=-84.47
CNS ~ Mass + Fat + HEART + MUSCLE
```

	Df	Sum of Sq	RSS	AIC
- Fat	1	0.04057	3.5007	-86.014
- Mass	1	0.09476	3.5549	-85.415
<none>			3.4601	-84.468
- HEART	1	0.40303	3.8631	-82.171
- MUSCLE	1	0.41063	3.8707	-82.095

```
Step: AIC=-86.01
CNS ~ Mass + HEART + MUSCLE
```

	Df	Sum of Sq	RSS	AIC
<none>			3.5007	-86.014
- Mass	1	0.35155	3.8522	-84.282
- HEART	1	0.37281	3.8735	-84.067
- MUSCLE	1	0.85479	4.3555	-79.493

```
summary(m9)
```

```
Call:
lm(formula = CNS ~ Mass + HEART + MUSCLE, data = data_1)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-0.74267 -0.11882 -0.00818  0.10790  0.84702
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.1840     0.8704  -0.211  0.83383
Mass        -0.8197     0.4372  -1.875  0.06919 .
HEART        0.3855     0.1997   1.931  0.06166 .
MUSCLE       1.2436     0.4254   2.923  0.00603 **
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.3163 on 35 degrees of freedom
Multiple R-squared:  0.9723,    Adjusted R-squared:  0.97
F-statistic: 410.2 on 3 and 35 DF,  p-value: < 2.2e-16
```

And both:

```
# Stepwise model selection - both  
m10 = step(m7, direction = 'both')
```

Start: AIC=-80.52

CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE

	Df	Sum of Sq	RSS	AIC
- FFWT	1	0.00001	3.4552	-82.523
- BONE	1	0.00422	3.4594	-82.476
- Mass	1	0.00517	3.4604	-82.465
- Fat	1	0.02092	3.4761	-82.288
<none>			3.4552	-80.523
- MUSCLE	1	0.30975	3.7650	-79.175
- HEART	1	0.39468	3.8499	-78.305

Step: AIC=-82.52

CNS ~ Mass + Fat + HEART + MUSCLE + BONE

	Df	Sum of Sq	RSS	AIC
- BONE	1	0.00487	3.4601	-84.468
- Fat	1	0.04499	3.5002	-84.019
- Mass	1	0.05110	3.5063	-83.951
<none>			3.4552	-82.523
+ FFWT	1	0.00001	3.4552	-80.523
- MUSCLE	1	0.38148	3.8367	-80.439
- HEART	1	0.39621	3.8514	-80.289

Step: AIC=-84.47

CNS ~ Mass + Fat + HEART + MUSCLE

	Df	Sum of Sq	RSS	AIC
- Fat	1	0.04057	3.5007	-86.014
- Mass	1	0.09476	3.5549	-85.415
<none>			3.4601	-84.468
+ BONE	1	0.00487	3.4552	-82.523
+ FFWT	1	0.00067	3.4594	-82.476
- HEART	1	0.40303	3.8631	-82.171
- MUSCLE	1	0.41063	3.8707	-82.095

Step: AIC=-86.01

```
CNS ~ Mass + HEART + MUSCLE
```

	Df	Sum of Sq	RSS	AIC
<none>			3.5007	-86.014
+ Fat	1	0.04057	3.4601	-84.468
- Mass	1	0.35155	3.8522	-84.282
+ FFWT	1	0.01839	3.4823	-84.219
- HEART	1	0.37281	3.8735	-84.067
+ BONE	1	0.00046	3.5002	-84.019
- MUSCLE	1	0.85479	4.3555	-79.493

```
summary(m10)
```

Call:

```
lm(formula = CNS ~ Mass + HEART + MUSCLE, data = data_1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.74267	-0.11882	-0.00818	0.10790	0.84702

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.1840	0.8704	-0.211	0.83383
Mass	-0.8197	0.4372	-1.875	0.06919 .
HEART	0.3855	0.1997	1.931	0.06166 .
MUSCLE	1.2436	0.4254	2.923	0.00603 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3163 on 35 degrees of freedom

Multiple R-squared: 0.9723, Adjusted R-squared: 0.97

F-statistic: 410.2 on 3 and 35 DF, p-value: < 2.2e-16

If you want to get fancy, we can even look at every possible model

```
library(MuMIn)
```

```
# Set global options to avoid error
options('na.action' = na.fail)
```

```
# Run dredge to get full selection table
dredge(m7, rank = 'AIC')
```

Fixed term is "(Intercept)"

```
Global model call: lm(formula = CNS ~ Mass + Fat + FFWT + HEART + MUSCLE + BONE,
  data = data_1)
```

Model selection table

	(Intrc)	BONE	Fat	FFWT	HEART	Mass	MUSCL	df	logLik	AIC
57	-0.1840				0.3855	-0.8197	1.2440	5	-8.332	26.7
43	-1.1710	-0.124700			0.3924		0.5758	5	-8.631	27.3
45	-0.2475		-0.85860		0.3612		1.2890	5	-8.961	27.9
59	-0.4891	-0.061790			0.4060	-0.5708	1.0530	6	-8.104	28.2
47	-0.5634	-0.096360	-0.57650		0.4140		1.0840	6	-8.124	28.2
41	-1.1430				0.2965		0.5107	4	-10.198	28.4
61	-0.3049			0.58870	0.3869	-1.2410	1.0870	6	-8.229	28.5
49	-1.7220					-0.6191	1.3760	4	-10.305	28.6
58	-0.2132	-0.01926			0.3888	-0.8022	1.2420	6	-8.329	28.7
33	-2.2060						0.7649	3	-11.347	28.7
44	-1.2470	-0.17370	-0.114900		0.4320		0.6975	6	-8.363	28.7
35	-2.4620		-0.085950				0.8664	4	-10.603	29.2
42	-1.2550	-0.24820			0.3639		0.6919	5	-9.670	29.3
37	-1.7080		-0.64360				1.3900	4	-10.670	29.3
46	-0.3146	-0.03584		-0.80970	0.3672		1.2710	6	-8.953	29.9
60	-0.6121	-0.06454	-0.066880		0.4188	-0.4915	1.0320	7	-8.077	30.2
63	-0.4991		-0.070960	-0.17140	0.4086	-0.4111	1.0710	7	-8.101	30.2
48	-0.6664	-0.05386	-0.097010	-0.50110	0.4234		1.0550	7	-8.106	30.2
15	-1.3260		-0.141700	0.53690	0.4770			5	-10.136	30.3
62	-0.4490	-0.07620		0.72770	0.4004	-1.2710	1.0430	7	-8.195	30.4
29	-0.9446			2.32300	0.4383	-1.8930		5	-10.222	30.4
53	-1.8390			0.54460		-1.0080	1.2310	5	-10.226	30.5
50	-1.4810	0.10920				-0.7284	1.3790	5	-10.226	30.5
34	-2.3590	-0.10740					0.8684	4	-11.244	30.5
51	-1.9150		-0.030810			-0.4897	1.2850	5	-10.253	30.5
27	-1.3680		-0.174800		0.4968	0.5506		5	-10.291	30.6
39	-2.0530		-0.062690	-0.43960			1.2660	5	-10.335	30.7
36	-2.5030	-0.03440	-0.083230				0.8963	5	-10.593	31.2
38	-1.4490	0.10160		-0.79230			1.4370	5	-10.612	31.2
16	-1.6470	-0.20880	-0.139600	0.71490	0.5072			6	-9.877	31.8
30	-1.3160	-0.23860		2.54100	0.4741	-1.9060		6	-9.885	31.8

13	-1.0500			0.42240	0.4110			4	-11.885	31.8
31	-1.1730		-0.087570	1.35400	0.4642	-0.8572		6	-10.045	32.1
64	-0.6114	-0.06339	-0.068180	-0.02606	0.4189	-0.4687	1.0350	8	-8.077	32.2
28	-1.6410	-0.16770	-0.182100		0.5254	0.6983		6	-10.119	32.2
54	-1.6320	0.07991		0.40060		-0.9854	1.2720	6	-10.189	32.4
52	-1.6650	0.09609	-0.024700			-0.6115	1.3050	6	-10.194	32.4
55	-1.8600		-0.005303	0.48770		-0.9454	1.2310	6	-10.225	32.5
40	-1.7890	0.10350	-0.062910	-0.59040			1.3130	6	-10.273	32.5
21	-2.8120			2.53800		-1.7250		4	-12.557	33.1
25	-0.7492				0.4796	0.3550		4	-12.592	33.2
14	-1.4090	-0.23040		0.62070	0.4454			5	-11.597	33.2
32	-1.4890	-0.22300	-0.075840	1.68700	0.4941	-1.0080		7	-9.751	33.5
5	-2.8020			0.78490				3	-13.802	33.6
7	-3.2380		-0.110900	0.92010				4	-12.820	33.6
56	-1.6680	0.08319	-0.011140	0.27490		-0.8521	1.2720	7	-10.186	34.4
19	-3.4510		-0.166800			0.9721		4	-13.253	34.5
9	0.9951				0.8820			3	-14.404	34.8
12	0.3485	0.33210	-0.111800		0.6667			5	-12.415	34.8
10	0.3284	0.25560			0.5985			4	-13.429	34.9
22	-3.0100	-0.08987		2.62700		-1.7240		5	-12.513	35.0
23	-2.8680		-0.014310	2.38200		-1.5540		5	-12.553	35.1
26	-0.8465	-0.06583			0.4905	0.4097		5	-12.568	35.1
6	-3.0010	-0.09063		0.87490				4	-13.760	35.5
8	-3.3620	-0.05836	-0.109700	0.97660				5	-12.802	35.6
11	1.1350		-0.071490		0.9798			4	-13.981	36.0
17	-2.7910					0.7710		3	-15.079	36.2
20	-3.4240	0.01165	-0.166400			0.9601		5	-13.252	36.5
24	-3.0350	-0.08761	-0.007846	2.53900		-1.6310		6	-12.511	37.0
18	-2.5850	0.09433				0.6790		4	-15.033	38.1
2	-1.0500	0.78560						3	-17.378	40.8
4	-1.1140	0.84970	-0.052420					4	-17.190	42.4
3	-0.1571		0.599000					3	-35.811	77.6
1	1.3230							2	-78.299	160.6
delta weight										
57	0.00	0.104								
43	0.60	0.077								
45	1.26	0.056								
59	1.55	0.048								
47	1.58	0.047								
41	1.73	0.044								
61	1.79	0.042								
49	1.95	0.039								
58	1.99	0.038								

33	2.03	0.038
44	2.06	0.037
35	2.54	0.029
42	2.68	0.027
37	2.68	0.027
46	3.24	0.021
60	3.49	0.018
63	3.54	0.018
48	3.55	0.018
15	3.61	0.017
62	3.73	0.016
29	3.78	0.016
53	3.79	0.016
50	3.79	0.016
34	3.82	0.015
51	3.84	0.015
27	3.92	0.015
39	4.01	0.014
36	4.52	0.011
38	4.56	0.011
16	5.09	0.008
30	5.11	0.008
13	5.11	0.008
31	5.43	0.007
64	5.49	0.007
28	5.57	0.006
54	5.71	0.006
52	5.72	0.006
55	5.79	0.006
40	5.88	0.005
21	6.45	0.004
25	6.52	0.004
14	6.53	0.004
32	6.84	0.003
5	6.94	0.003
7	6.98	0.003
56	7.71	0.002
19	7.84	0.002
9	8.14	0.002
12	8.17	0.002
10	8.19	0.002
22	8.36	0.002
23	8.44	0.002

26	8.47	0.002
6	8.86	0.001
8	8.94	0.001
11	9.30	0.001
17	9.49	0.001
20	9.84	0.001
24	10.36	0.001
18	11.40	0.000
2	14.09	0.000
4	15.72	0.000
3	50.96	0.000
1	133.93	0.000

Models ranked by AIC(x)

Here, we're ranking models by AIC. AIC balances fit with model complexity. Lower values of AIC are considered better. Generally, 2 is used as a rule of thumb for AIC - if delta AIC is >2 , the lower AIC model is considered better. If delta AIC is <2 , the support for the best model is weak, and the models could even be considered "tied".

5

5.1 Tips for your Assignment:

Some things you may want to think about for your assignment:

1. What role is collinearity playing in your assignment? Is it something you should be concerned about? Why or why not?
2. What does it mean if your assumptions are violated? How would you fix it? Is it worth fixing it? Why or why not?
3. How would you interpret your statistical results biologically? You don't have to be right, but don't be vague, and don't contradict your results.

6 Assignment 1e: Bayesian Data Analysis

Assignment 1e is an introduction Bayesian data analysis, using Bayesian general linear models.

For this tutorial, we'll be using `cuse.csv` .

6.1 Looking at the Data

```
# Load in data
data = read.csv('cuse.csv')

# Look at the data structure
head(data)
```

	X	age	education	wantsMore	notUsing	using
1	1	<25	low	yes	53	6
2	2	<25	low	no	10	4
3	3	<25	high	yes	212	52
4	4	<25	high	no	50	10
5	5	25-29	low	yes	60	14
6	6	25-29	low	no	19	10

```
dim(data)
```

```
[1] 16 6
```

Our data contains 16 observations of 5 variables - a binomial matrix of how many women are using or not using birth control within 16 groups, and three categorical predictors - age, expressed as a bin, education, and whether they want more children. The first column is a duplicate of our row names. We can get rid of that:

```
# Remove column 1
data = data[,-1]
head(data)
```

	age	education	wantsMore	notUsing	using
1	<25	low	yes	53	6
2	<25	low	no	10	4
3	<25	high	yes	212	52
4	<25	high	no	50	10
5	25-29	low	yes	60	14
6	25-29	low	no	19	10

6.2 Binomial GLM

Binomial general linear models are used to calculate the probability of a binomial response - in this case, whether someone is using or not using birth control. Binomial GLM response can be fed in either as a true/false set, or as a matrix of successes and failures. According to ?family, we need to feed in the data with successes first and failures second. Let's create the matrix:

```
# create response matrix
resp = cbind(data$using, data$notUsing)
head(resp)
```

	[,1]	[,2]
[1,]	6	53
[2,]	4	10
[3,]	52	212
[4,]	10	50
[5,]	14	60
[6,]	10	19

In this case, all of our variables are categorical, and they are currently stored as characters:

```
# Check predictor classes
class(data$age)
```

```
[1] "character"
```

```
class(data$education)
```

```
[1] "character"
```

```
class(data$wantsMore)
```

```
[1] "character"
```

These should function fine as categorical variables. Let's make our GLM:

```
# Run GLM
m1 = glm(resp ~ age + education + wantsMore, family = 'binomial', data = data)
summary(m1) # Summary
```

Call:

```
glm(formula = resp ~ age + education + wantsMore, family = "binomial",
     data = data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-0.8082	0.1590	-5.083	3.71e-07	***
age25-29	0.3894	0.1759	2.214	0.02681	*
age30-39	0.9086	0.1646	5.519	3.40e-08	***
age40-49	1.1892	0.2144	5.546	2.92e-08	***
educationlow	-0.3250	0.1240	-2.620	0.00879	**
wantsMoreyes	-0.8330	0.1175	-7.091	1.33e-12	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 165.772 on 15 degrees of freedom
Residual deviance: 29.917 on 10 degrees of freedom
AIC: 113.43

Number of Fisher Scoring iterations: 4

In our summary we see we have 5 predictors: The age bin, low education, and wanting more kids. High education and not wanting more kids are missing because these variables are binary, so we only need one variable to differentiate them. We can also see the values of our model coefficients, their standard errors, and the model AIC.

6.3 Making it Bayesian

The default GLM function is frequentist (that's why we have p-values). Now lets try a Bayesian approach:

```
# Stan  
library(rstanarm)
```

Loading required package: Rcpp

This is rstanarm version 2.32.1

- See <https://mc-stan.org/rstanarm/articles/priors> for changes to default priors!
- Default priors may change, so it's safest to specify priors, even if equivalent to the default
- For execution on a local, multicore CPU with excess RAM we recommend calling
`options(mc.cores = parallel::detectCores())`

```
library(bayesplot)
```

This is bayesplot version 1.11.1

- Online documentation and vignettes at mc-stan.org/bayesplot
- bayesplot theme set to `bayesplot::theme_default()`
 - * Does `_not_` affect other ggplot2 plots
 - * See `?bayesplot_theme_set` for details on theme setting


```
library(shinystan)
```

Loading required package: shiny

This is shinystan version 2.6.0

```
library(ggplot2)
```

```
# Run glm
```

```
m2 = stan_glm(resp ~ age + education + wantsMore, family = 'binomial', data = data)
```

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 6.6e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.66 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [20%] (Warmup)

Chain 1: Iteration: 600 / 2000 [30%] (Warmup)

Chain 1: Iteration: 800 / 2000 [40%] (Warmup)

Chain 1: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 1: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 1: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 1: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 1: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 1: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.164 seconds (Warm-up)

Chain 1: 0.166 seconds (Sampling)

Chain 1: 0.33 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 2.2e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.22 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration: 1 / 2000 [0%] (Warmup)
Chain 2: Iteration: 200 / 2000 [10%] (Warmup)
Chain 2: Iteration: 400 / 2000 [20%] (Warmup)
Chain 2: Iteration: 600 / 2000 [30%] (Warmup)
Chain 2: Iteration: 800 / 2000 [40%] (Warmup)
Chain 2: Iteration: 1000 / 2000 [50%] (Warmup)
Chain 2: Iteration: 1001 / 2000 [50%] (Sampling)
Chain 2: Iteration: 1200 / 2000 [60%] (Sampling)
Chain 2: Iteration: 1400 / 2000 [70%] (Sampling)
Chain 2: Iteration: 1600 / 2000 [80%] (Sampling)
Chain 2: Iteration: 1800 / 2000 [90%] (Sampling)
Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.167 seconds (Warm-up)
Chain 2: 0.167 seconds (Sampling)
Chain 2: 0.334 seconds (Total)
Chain 2:

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 3).

Chain 3:
Chain 3: Gradient evaluation took 2.4e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.24 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration: 1 / 2000 [0%] (Warmup)
Chain 3: Iteration: 200 / 2000 [10%] (Warmup)
Chain 3: Iteration: 400 / 2000 [20%] (Warmup)
Chain 3: Iteration: 600 / 2000 [30%] (Warmup)
Chain 3: Iteration: 800 / 2000 [40%] (Warmup)
Chain 3: Iteration: 1000 / 2000 [50%] (Warmup)
Chain 3: Iteration: 1001 / 2000 [50%] (Sampling)
Chain 3: Iteration: 1200 / 2000 [60%] (Sampling)
Chain 3: Iteration: 1400 / 2000 [70%] (Sampling)
Chain 3: Iteration: 1600 / 2000 [80%] (Sampling)
Chain 3: Iteration: 1800 / 2000 [90%] (Sampling)
Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 3:

```
Chain 3: Elapsed Time: 0.164 seconds (Warm-up)
Chain 3:           0.166 seconds (Sampling)
Chain 3:           0.33 seconds (Total)
Chain 3:
```

```
SAMPLING FOR MODEL 'binomial' NOW (CHAIN 4).
```

```
Chain 4:
Chain 4: Gradient evaluation took 2e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.2 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
Chain 4: Iteration:  1001 / 2000 [ 50%] (Sampling)
Chain 4: Iteration:  1200 / 2000 [ 60%] (Sampling)
Chain 4: Iteration:  1400 / 2000 [ 70%] (Sampling)
Chain 4: Iteration:  1600 / 2000 [ 80%] (Sampling)
Chain 4: Iteration:  1800 / 2000 [ 90%] (Sampling)
Chain 4: Iteration:  2000 / 2000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 0.167 seconds (Warm-up)
Chain 4:           0.155 seconds (Sampling)
Chain 4:           0.322 seconds (Total)
Chain 4:
```

```
summary(m2) # Summary
```

```
Model Info:
```

```
function:    stan_glm
family:      binomial [logit]
formula:     resp ~ age + education + wantsMore
algorithm:   sampling
sample:      4000 (posterior sample size)
priors:      see help('prior_summary')
observations: 16
```

```
predictors:    6
```

Estimates:

	mean	sd	10%	50%	90%
(Intercept)	-0.8	0.2	-1.0	-0.8	-0.6
age25-29	0.4	0.2	0.2	0.4	0.6
age30-39	0.9	0.2	0.7	0.9	1.1
age40-49	1.2	0.2	0.9	1.2	1.5
educationlow	-0.3	0.1	-0.5	-0.3	-0.2
wantsMoreyes	-0.8	0.1	-1.0	-0.8	-0.7

Fit Diagnostics:

	mean	sd	10%	50%	90%
mean_PPD	31.7	1.6	29.7	31.7	33.7

The mean_ppd is the sample average posterior predictive distribution of the outcome variable


MCMC diagnostics

	mcse	Rhat	n_eff
(Intercept)	0.0	1.0	2138
age25-29	0.0	1.0	2362
age30-39	0.0	1.0	2051
age40-49	0.0	1.0	2412
educationlow	0.0	1.0	2772
wantsMoreyes	0.0	1.0	3022
mean_PPD	0.0	1.0	3638
log-posterior	0.0	1.0	1812

For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective

The `stan_glm` function automatically feeds our model into Stan, which is a Hamiltonian Markov Chain Monte Carlo (MCMC) sampler. Running `summary` on our model gives us some model diagnostics - all our Rhat values are 1 and all our n_eff values are well into the thousands, both of which are a good sign. We can also do some visual checks and tests:


```
# Trace plot
plot(m2, 'trace')
```



These are trace plots, which show us the parameter values selected for each iteration of the MCMC chain. We want these to look “fuzzy” - that indicates the sampler is exploring the full range of possible values. If these lines were to be flat, that would indicate the sampler got “stuck” and didn’t sample the full posterior distributions. These look good.


Lets look at our posteriors:

```
# Plot parameter values with uncertainties
plot(m2, prob_outer = 0.95)
```



```
# Plot posterior distributions  
plot(m2, 'mcmc_hist')
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



These plots both give us an idea of our parameter values and their posterior distributions. The former plot shows the median parameter estimates (circle), their 50% quantiles (dark blue box), and their 95% quantiles (thin blue line). The latter shows histograms of the posterior distributions of each of our parameters.

We can also pull out our coefficients and posteriors directly


```
# Model coefficients
m2$coefficients
```

```
(Intercept)    age25-29    age30-39    age40-49 educationlow wantsMoreyes
-0.8151250     0.3946477     0.9131451     1.1978136    -0.3266466    -0.8327661
```

```
# Model posteriors
posterior <- as.matrix(m2)

# Plot model posteriors (95% quantile)
plot_title <- ggtitle("Posterior distributions with medians and 95% credible intervals")

mcmc_areas(posterior, pars = names(m2$coefficients),
            prob = 0.95) + plot_title
```



How would you interpret these plots?

6.4 Adding Priors

Lets try adding some priors:

```
# Run glm with priors
m3 = stan_glm(resp ~ age + education + wantsMore, family = 'binomial', data = data,
              prior = normal(location = c(0.2, 1.5, 2, -1, -0.25), # Normal priors, means
                             scale = c(0.03, 0.03, 0.03, 0.03, 0.03))) # And standard deviation
```

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 3.8e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.38 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)


```

Chain 1: Iteration: 200 / 2000 [ 10%] (Warmup)
Chain 1: Iteration: 400 / 2000 [ 20%] (Warmup)
Chain 1: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 1: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 1:
Chain 1: Elapsed Time: 0.215 seconds (Warm-up)
Chain 1:           0.12 seconds (Sampling)
Chain 1:           0.335 seconds (Total)
Chain 1:

```

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 2).

```

Chain 2:
Chain 2: Gradient evaluation took 2.5e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.25 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.202 seconds (Warm-up)
Chain 2:           0.118 seconds (Sampling)
Chain 2:           0.32 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 1.7e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.17 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:    1 / 2000 [  0%]  (Warmup)
Chain 3: Iteration:   200 / 2000 [ 10%]  (Warmup)
Chain 3: Iteration:   400 / 2000 [ 20%]  (Warmup)
Chain 3: Iteration:   600 / 2000 [ 30%]  (Warmup)
Chain 3: Iteration:   800 / 2000 [ 40%]  (Warmup)
Chain 3: Iteration:  1000 / 2000 [ 50%]  (Warmup)
Chain 3: Iteration:  1001 / 2000 [ 50%]  (Sampling)
Chain 3: Iteration:  1200 / 2000 [ 60%]  (Sampling)
Chain 3: Iteration:  1400 / 2000 [ 70%]  (Sampling)
Chain 3: Iteration:  1600 / 2000 [ 80%]  (Sampling)
Chain 3: Iteration:  1800 / 2000 [ 90%]  (Sampling)
Chain 3: Iteration:  2000 / 2000 [100%]  (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.196 seconds (Warm-up)
Chain 3:                  0.115 seconds (Sampling)
Chain 3:                  0.311 seconds (Total)
Chain 3:

```

SAMPLING FOR MODEL 'binomial' NOW (CHAIN 4).

```

Chain 4:
Chain 4: Gradient evaluation took 1.9e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.19 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration:    1 / 2000 [  0%]  (Warmup)
Chain 4: Iteration:   200 / 2000 [ 10%]  (Warmup)
Chain 4: Iteration:   400 / 2000 [ 20%]  (Warmup)
Chain 4: Iteration:   600 / 2000 [ 30%]  (Warmup)
Chain 4: Iteration:   800 / 2000 [ 40%]  (Warmup)
Chain 4: Iteration:  1000 / 2000 [ 50%]  (Warmup)
Chain 4: Iteration:  1001 / 2000 [ 50%]  (Sampling)
Chain 4: Iteration:  1200 / 2000 [ 60%]  (Sampling)
Chain 4: Iteration:  1400 / 2000 [ 70%]  (Sampling)
Chain 4: Iteration:  1600 / 2000 [ 80%]  (Sampling)
Chain 4: Iteration:  1800 / 2000 [ 90%]  (Sampling)
Chain 4: Iteration:  2000 / 2000 [100%]  (Sampling)

```

```
Chain 4:
Chain 4: Elapsed Time: 0.245 seconds (Warm-up)
Chain 4:           0.112 seconds (Sampling)
Chain 4:           0.357 seconds (Total)
Chain 4:
```

```
summary(m3) # Summary
```

Model Info:

```
function:      stan_glm
family:        binomial [logit]
formula:       resp ~ age + education + wantsMore
algorithm:     sampling
sample:        4000 (posterior sample size)
priors:         see help('prior_summary')
observations:  16
predictors:    6
```

Estimates:

	mean	sd	10%	50%	90%
(Intercept)	-1.2	0.1	-1.3	-1.2	-1.1
age25-29	0.2	0.0	0.2	0.2	0.3
age30-39	1.5	0.0	1.4	1.5	1.5
age40-49	2.0	0.0	2.0	2.0	2.0
educationlow	-1.0	0.0	-1.0	-1.0	-0.9
wantsMoreyes	-0.3	0.0	-0.3	-0.3	-0.2

Fit Diagnostics:

	mean	sd	10%	50%	90%
mean_PPD	31.7	1.5	29.8	31.7	33.6

The mean_ppd is the sample average posterior predictive distribution of the outcome variable

MCMC diagnostics

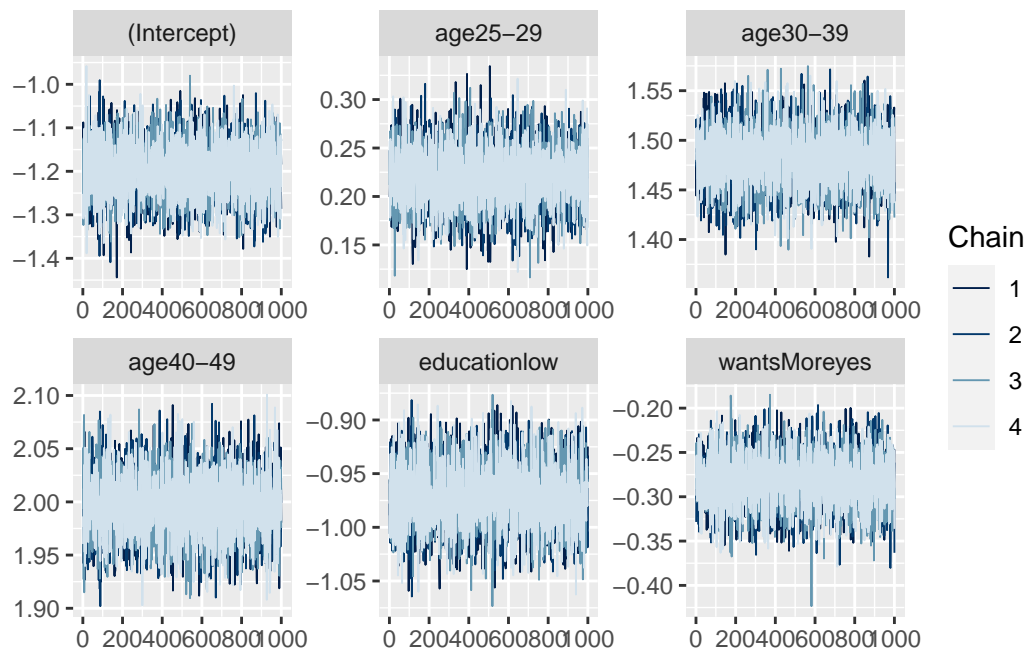
	mcse	Rhat	n_eff
(Intercept)	0.0	1.0	5041
age25-29	0.0	1.0	5661
age30-39	0.0	1.0	4981
age40-49	0.0	1.0	5371
educationlow	0.0	1.0	5241

```
wantsMoreyes 0.0 1.0 4932
mean_PPD     0.0 1.0 4974
log-posterior 0.0 1.0 1719
```


For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective

Lets look at our plots again:

```
# Trace plot
plot(m3, 'trace')
```




```
# Plot parameter values with uncertainties
plot(m3, prob_outer = 0.95)
```



```
# Plot posterior distributions  
plot(m3, 'mcmc_hist')
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.




```
# Model coefficients
m3$coefficients
```

(Intercept)	age25-29	age30-39	age40-49	educationlow	wantsMoreyes
-1.1993901	0.2201646	1.4818484	1.9952018	-0.9710766	-0.2824740

```
# Model posteriors
posterior <- as.matrix(m3)

# Plot model posteriors (95% quantile)
plot_title <- ggtitle("Posterior distributions with medians and 95% credible intervals")

mcmc_areas(posterior, pars = names(m3$coefficients),
            prob = 0.95) + plot_title
```



You can also look at all of your Stan model results using `shinystan` by running `launch_shinystan(model)`. Try it out on your end (it doesn't work in markdown)

6.5 Tips for your Assignment:

Some things you may want to think about for your assignment:

1. How do the results of these three models differ? Why or why not?
2. Do you believe certain models are more or less correct? Why or why not?
3. How would you interpret your statistical results biologically? You don't have to be right, but don't be vague, and don't contradict your results.